

# STAT5003

Week 4: High-dimensional visualisation and analytics in R

**Jaslene Lin**

*The University of Sydney*



THE UNIVERSITY OF  
**SYDNEY**

# Readings and R functions covered

## ! Important

- **Introduction to Statistical Learning**
- PCA Dimension reduction, see Section 10.2 - Clustering, see Section 10.3
- How to use *t*-SNE effectively
- **The Elements of Statistical Learning**
  - ➡ MDS, see Section 14.8
- **R** functions
- `hclust` (hierarchical clustering)
- `kmeans` (K-means clustering)
- `dist` (distance matrix)
- `Rtsne:::rtsne` (t-SNE)
- `cmdscale` (Multi-Dimensional Scaling)



# Unsupervised Learning

**Supervised Learning:** Given a dataset with features/predictors ( $X_{i1}, X_{i2}, \dots, X_{ip}$ ) and a target outcome variable ( $Y_i$ ), the goal is to learn a **model**  $f(X_{i1} \dots X_{ip})$  to explain/predict the target.

- example: Is there a relationship between BuildingArea and Price of houses in St Kilda?

**Unsupervised learning:** Given a dataset with features ( $X_{i1}, X_{i2}, \dots, X_{ip}$ ), the goal is to visualise, find patterns, subgroups/clusters within the data.

- example: Can we identify different types of houses in St Kilda based on their features?

# Clustering

# Clustering Overview

- Groups similar observations based on **a similarity measure**

## Clustering Goals

- **Compact Clusters:** Observations within a cluster should be **close together**.
- **Well-Separated Clusters:** Observations from different clusters should be **far apart**.
- Example algorithms:
  - ➡ **Hierarchical clustering**
  - ➡  **$K$ -means clustering**
  - ➡ Gaussian mixture model

# Typical methods

- Partitioning
  - ➡ Pre-specified number  $K$  of mutually exclusive and exhaustive groups
  - ➡ Iterate until criteria is met
- Hierarchical methods:
  - ➡ Agglomerative: Bottom up, more popular
  - ➡ Divisive: Top down, less popular
  - ➡ Display results with dendrogram

# *K*-Means Clustering

## Measuring Cluster Compactness

- For cluster  $C_k$ , we can define within-group sum of squares as:

$$\text{WSS}_k = \frac{1}{|C_k|} \sum_{i,j \in C_k} \|x_i - x_j\|^2, \quad k = 1, \dots, K$$

- This is the sum of all the pairwise squared Euclidean distances between observations in the  $k^{\text{th}}$  cluster, divided by total number of observations in the  $k^{\text{th}}$  cluster
- $\sum_{i,j \in C_k} \|x_i - x_j\|^2$  is the sum of all the pairwise squared Euclidean distances between observations in the  $k^{\text{th}}$  cluster.
- $|C_k|$  is the total number of observations in the  $k^{\text{th}}$  cluster

# *K*-Means Clustering

## Clustering Objective

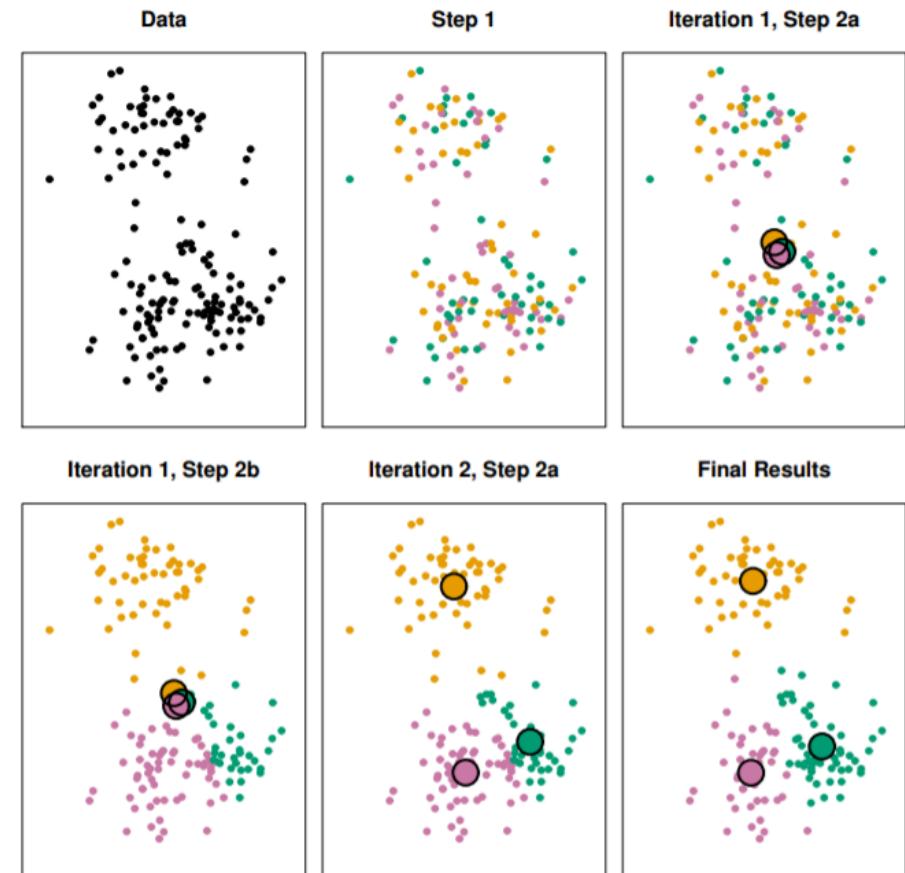
- To **minimize** the total within-group sum of squares criterion

$$\text{WSS}_{\text{Total}} = \sum_{k=1}^K \text{WSS}_k$$

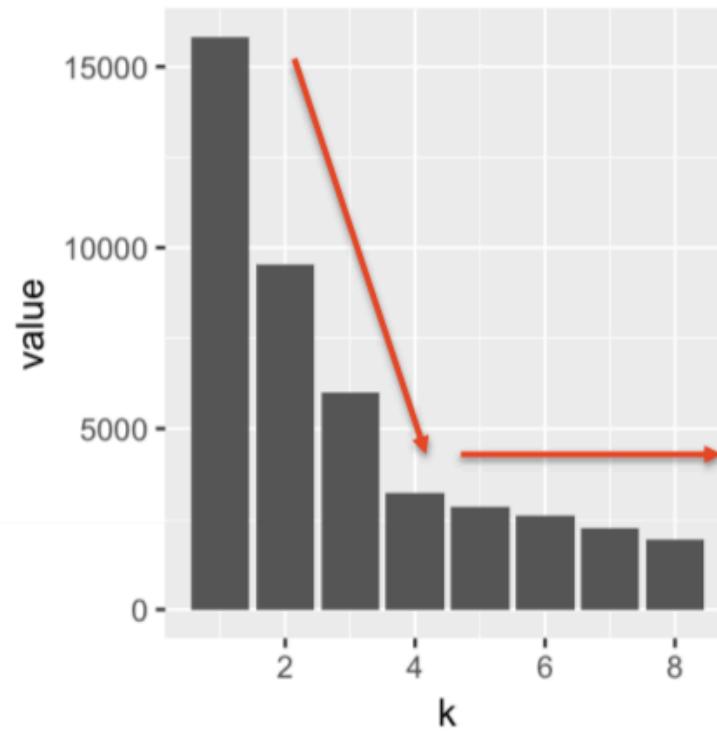
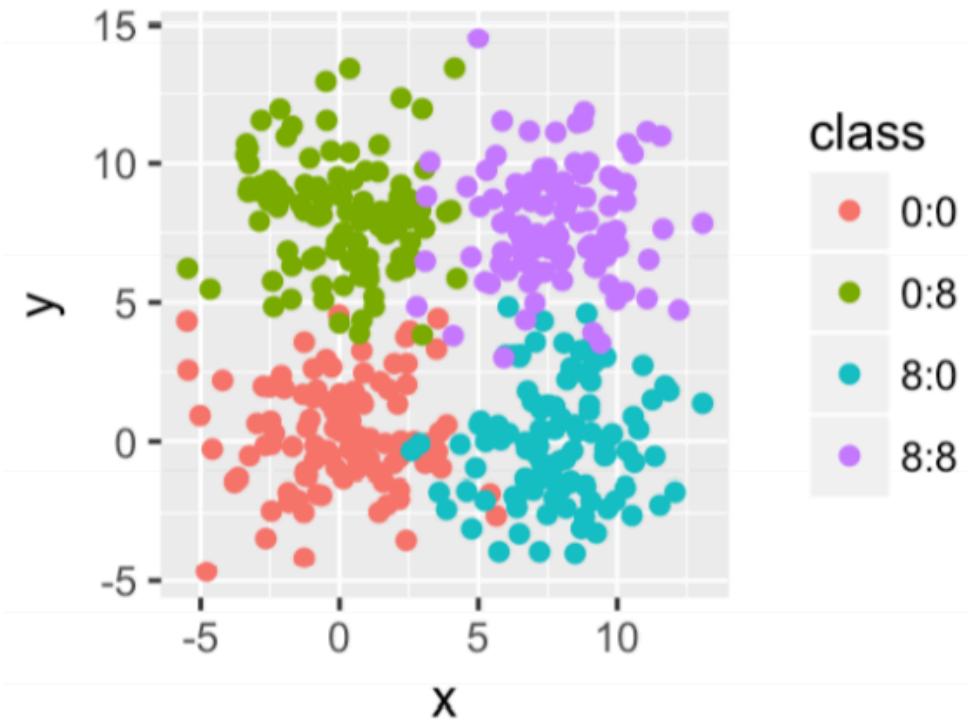
- The total within sum of square criterion will decrease as  $K$  increases
- Rule of thumb: Look for the elbow

# $K$ -Means Clustering

- Initialise each observation at random to a cluster.
- Iterate the following until the assignments stop changing:
  1. For each of the  $K$  clusters, compute the cluster centroid.
  2. Assign each observation to the cluster whose centroid is closest (where the closest is defined using the Euclidean distance)



# Choosing $K$ via Elbow plot



# *K*-Means properties

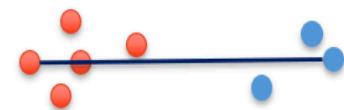
- The number of clusters  $K$  needs to be specified.
- Local solution and not necessarily global solution.
- Depends on starting values (the random starting values)-run the algorithms multiple times.
- Best for *compact, spherical* clusters.
- Does not work well when cluster sizes are different.

# Hierarchical clustering

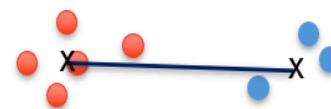
- Begin with every observation representing a single cluster
- At each iteration, merge the two closest clusters into one cluster
  - ➡ Needs a measure of **similarity/dissimilarity between two clusters**
  - ➡ These measures are called **linkages**.
- Linkages: Measure of dissimilarity between two sets of objects that determine how two set of objects are merged.
  - ➡ Single linkage
  - ➡ Complete linkage
  - ➡ Average Linkage



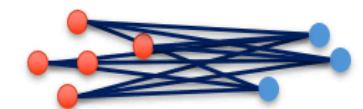
Single (minimum)



Complete (maximum)

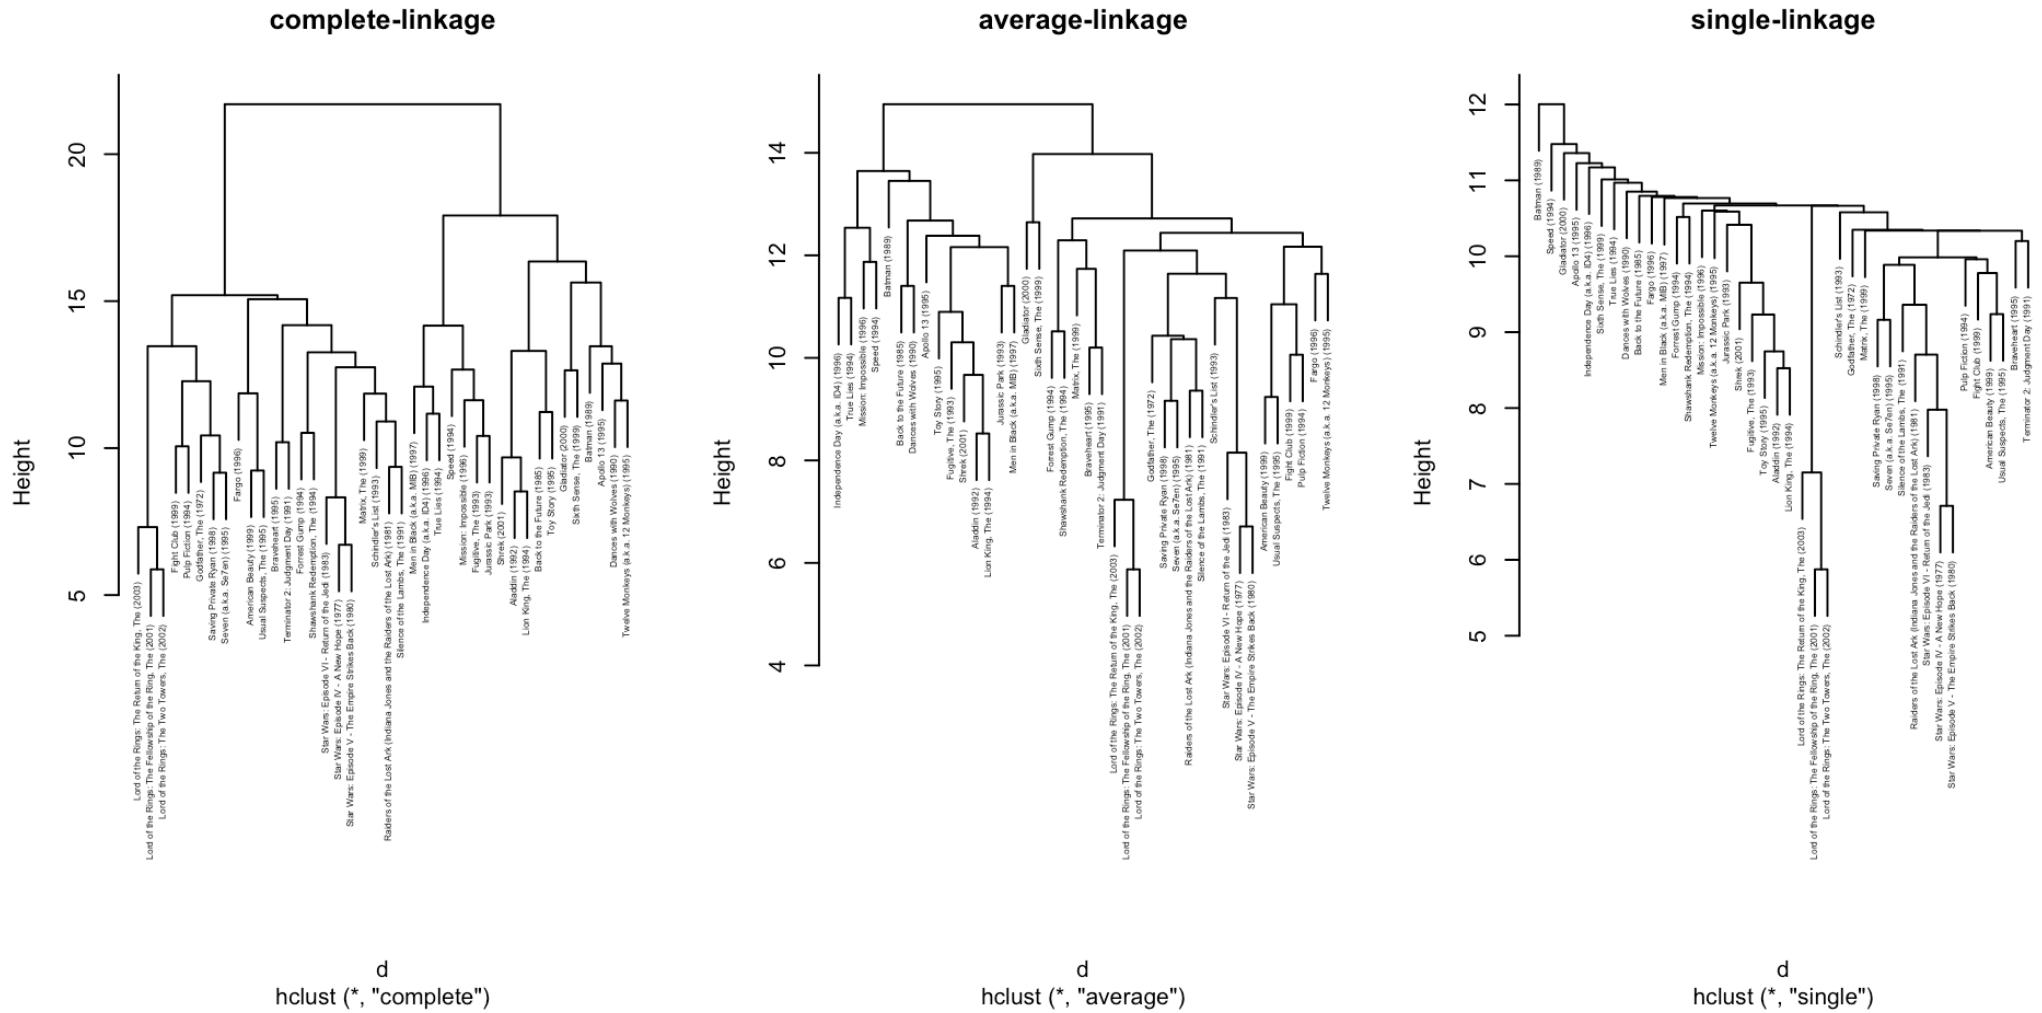


Distance between centroids



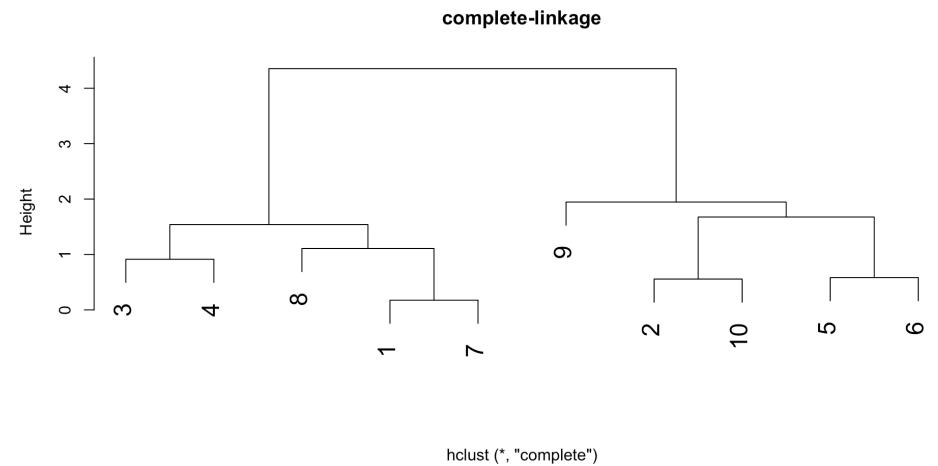
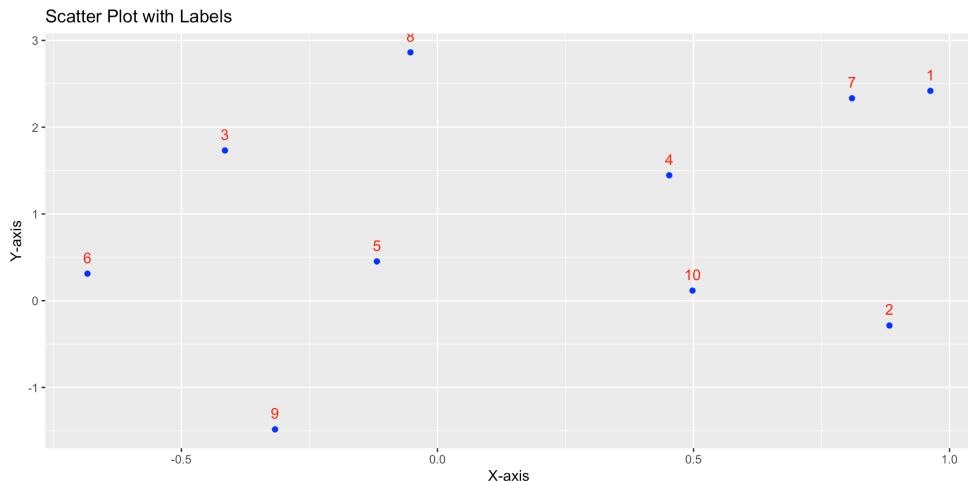
Average (mean) linkage

# Hierarchical clustering



# Hierarchical Clustering

## Dendrogram



# Dimension reduction: Principal Components Analysis (PCA)

# High dimensional data

- High-dimensional data refers to data set with **more features  $p$**  than observations  **$n$** 
  - ➡ Examples: in genetic data, we can easily measure 500k individual DNA mutations (human genome have ~3 billion base pairs of DNA), but experiments generally have < **1000** people, e.g.,  **$p \sim 500k$ ,  $n \sim 1000$**
- It is very hard to visualize high-dimensional data
  - ➡ Only have 2 (sometimes 3 or 4) dimensional canvas to create plots
- Many algorithms and methods have been designed for low dimensional data and would not work well for high-dimensional data
- To build a linear regression model data with **500k** features will result in **500k** parameters. This problem is underdetermined if we only have 1000 observations

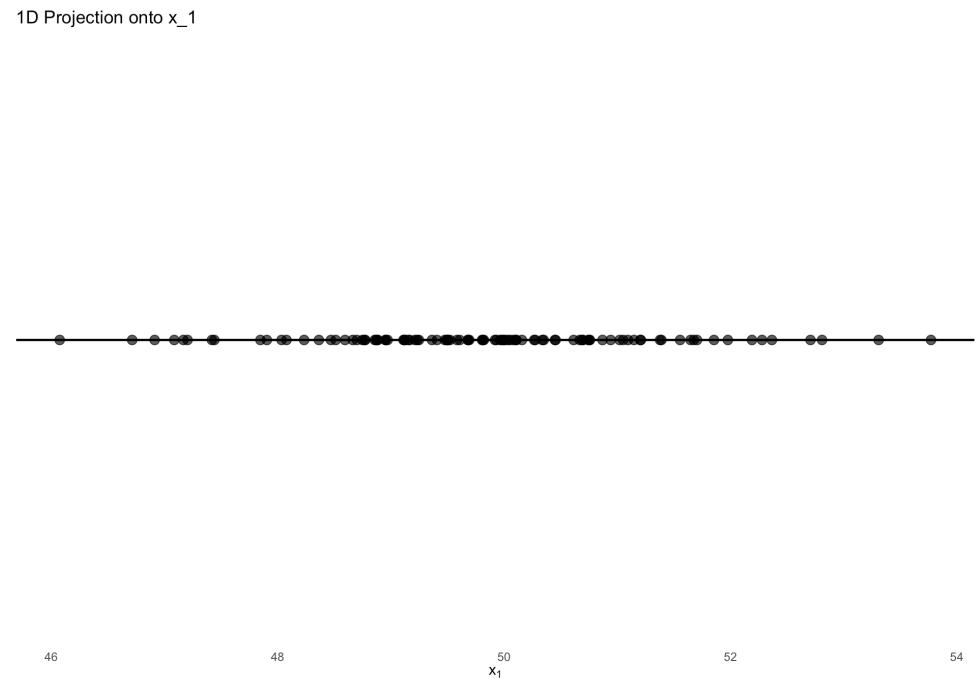
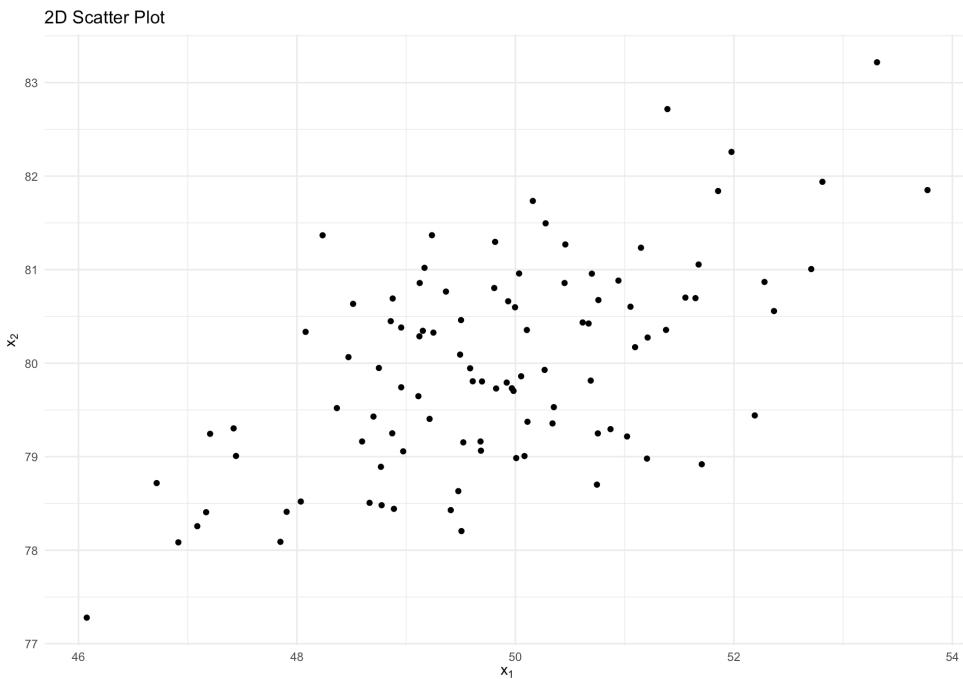
# Dimension reduction strategies

- Eliminate or remove features
  - ➡ Need to decide which features to be eliminated? Keep ones with high variance?
- Select features
  - ➡ Example: Lasso and ridge regression (coming soon in later module)
- Build or construct new features from existing ones
  - ➡ Replace many existing features with a single one
  - ➡ PCA and  $t$ -SNE

# PCA

- Suppose we have a data matrix  $\mathbf{X}$  with  $n$  observations and  $p$  features
  - ➡ Can we plot the data in a 2-dimensional plot?
- Naively, we can do all pairwise combinations, i.e., 1 vs 2, 1 vs 3,  $\dots$ , ( $p$  vs  $(p - 1)$ )
  - ➡  $\binom{p}{2} = \frac{p(p-1)}{2} = \mathcal{O}(p^2)$  different plots!
- Principal Component Analysis (PCA) helps us reorganise/transform the data into a new coordinate system.
  - ➡ The data still has  $p$  dimensions, but in a new set of axes.
  - ➡ The goal is to capture most of the variation using just the first few dimensions, making the data easier to analyse.

# Best way to represent 2d in 1d?



- Could use a single variable?  $x_1$  say?
- Or could remap  $x_1$  and  $x_2$  to a single variable
  - ➡ 
$$z = \phi_1 x_1 + \phi_2 x_2$$
- Can generalize this to many dimensions
  - ➡ 
$$z = \sum_{j=1}^p \phi_j x_j$$
- Pick the transformation that **maximises the variance!**

# Principal components

Start with a data matrix  $\mathbf{X}$ , and assume it has **mean zero**

$$\mathbf{X} = (X_1, X_2, \dots, X_p)$$

The first principal component is the **normalised linear combination of the features** that **maximises the variance** in the new component

$$Z_1 = \phi_{11}X_1 + \phi_{21}X_2 + \dots + \phi_{p1}X_p = \sum_{j=1}^p \phi_{j1}X_j = \mathbf{X}\boldsymbol{\phi}_1, \quad \boldsymbol{\phi}_1 = (\phi_{11}, \dots, \phi_{p1})^T$$

The elements  $\phi_{j1}$  are known as the loadings of the first principal component

- By normalised, we mean the squared loadings have to sum to 1, i.e.  $\sum_{j=1}^p \phi_{j1}^2 = 1 \Leftrightarrow \boldsymbol{\phi}_1^T \boldsymbol{\phi}_1 = 1$
- It is desired to maximise the sample variance of  $Z_{i1}, \dots, Z_{n1}$ , where  $Z_{i1}$  is the  $i$ th component of  $\mathbf{Z}_1$ :

$$\frac{1}{n-1} \sum_{i=1}^n \left( \sum_{j=1}^p \phi_{j1}x_{ij} \right)^2$$

## Principal component scores

- Given the principal component loadings, we can project our data matrix  $\mathbf{X}$  onto the principal component space  
➡ The projection is a linear combination of the sample feature values:

$$z_{i1} = \phi_{11}x_{i1} + \phi_{21}x_{i2} + \cdots + \phi_{p1}x_{ip}$$

This is known as the principal component score

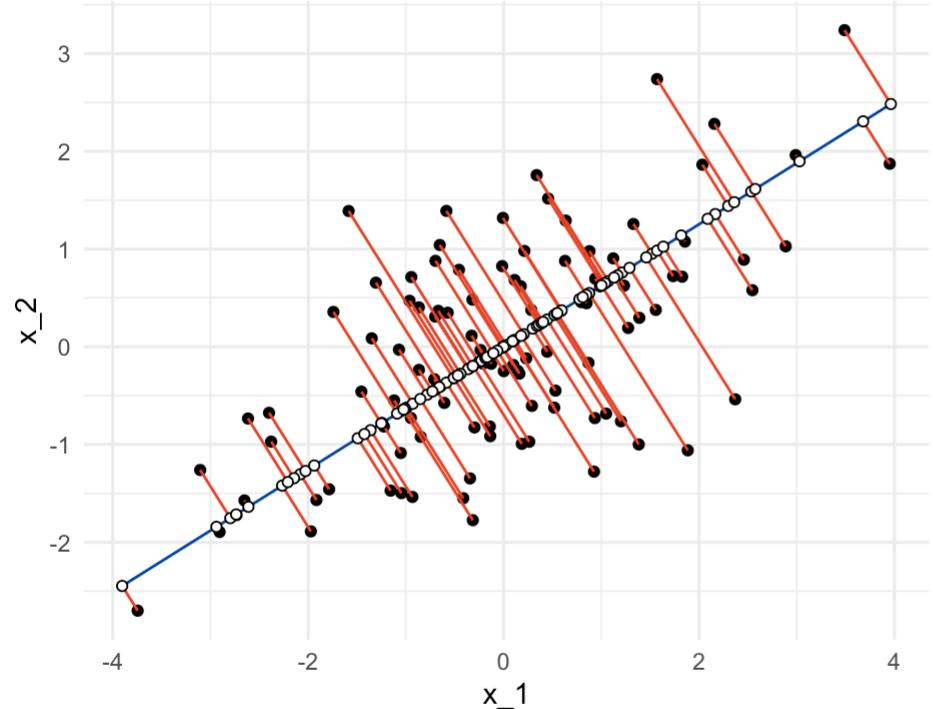
- The first principal component score vector is

$$\mathbf{Z}_1 = (z_{11}, z_{21}, \dots, z_{n1})$$

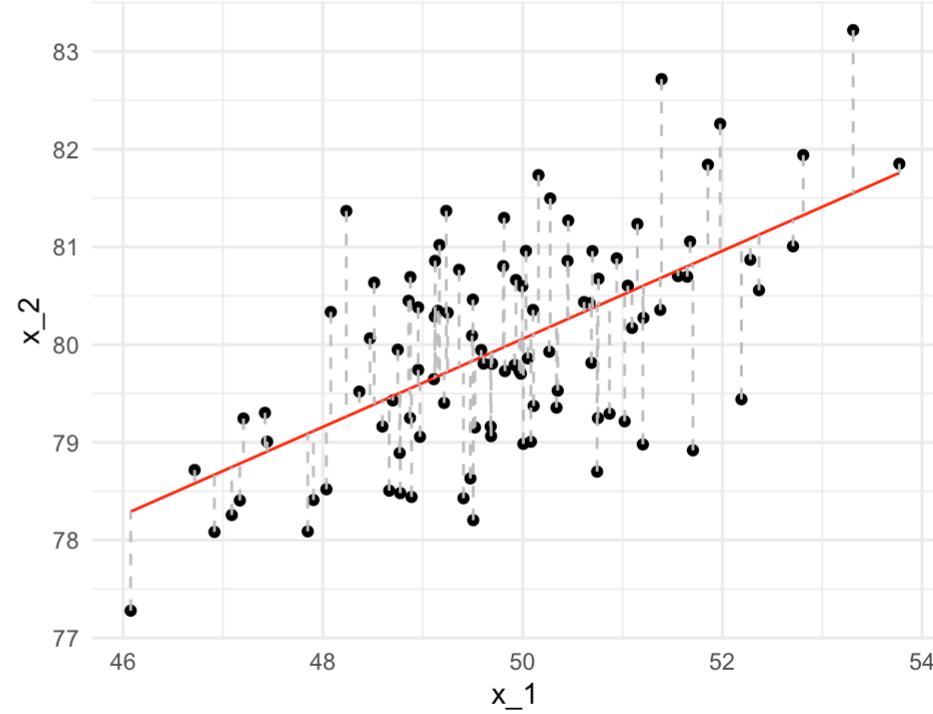
The principal component score vectors are uncorrelated (orthogonal)

# Geometric interpretation

First Principle Component

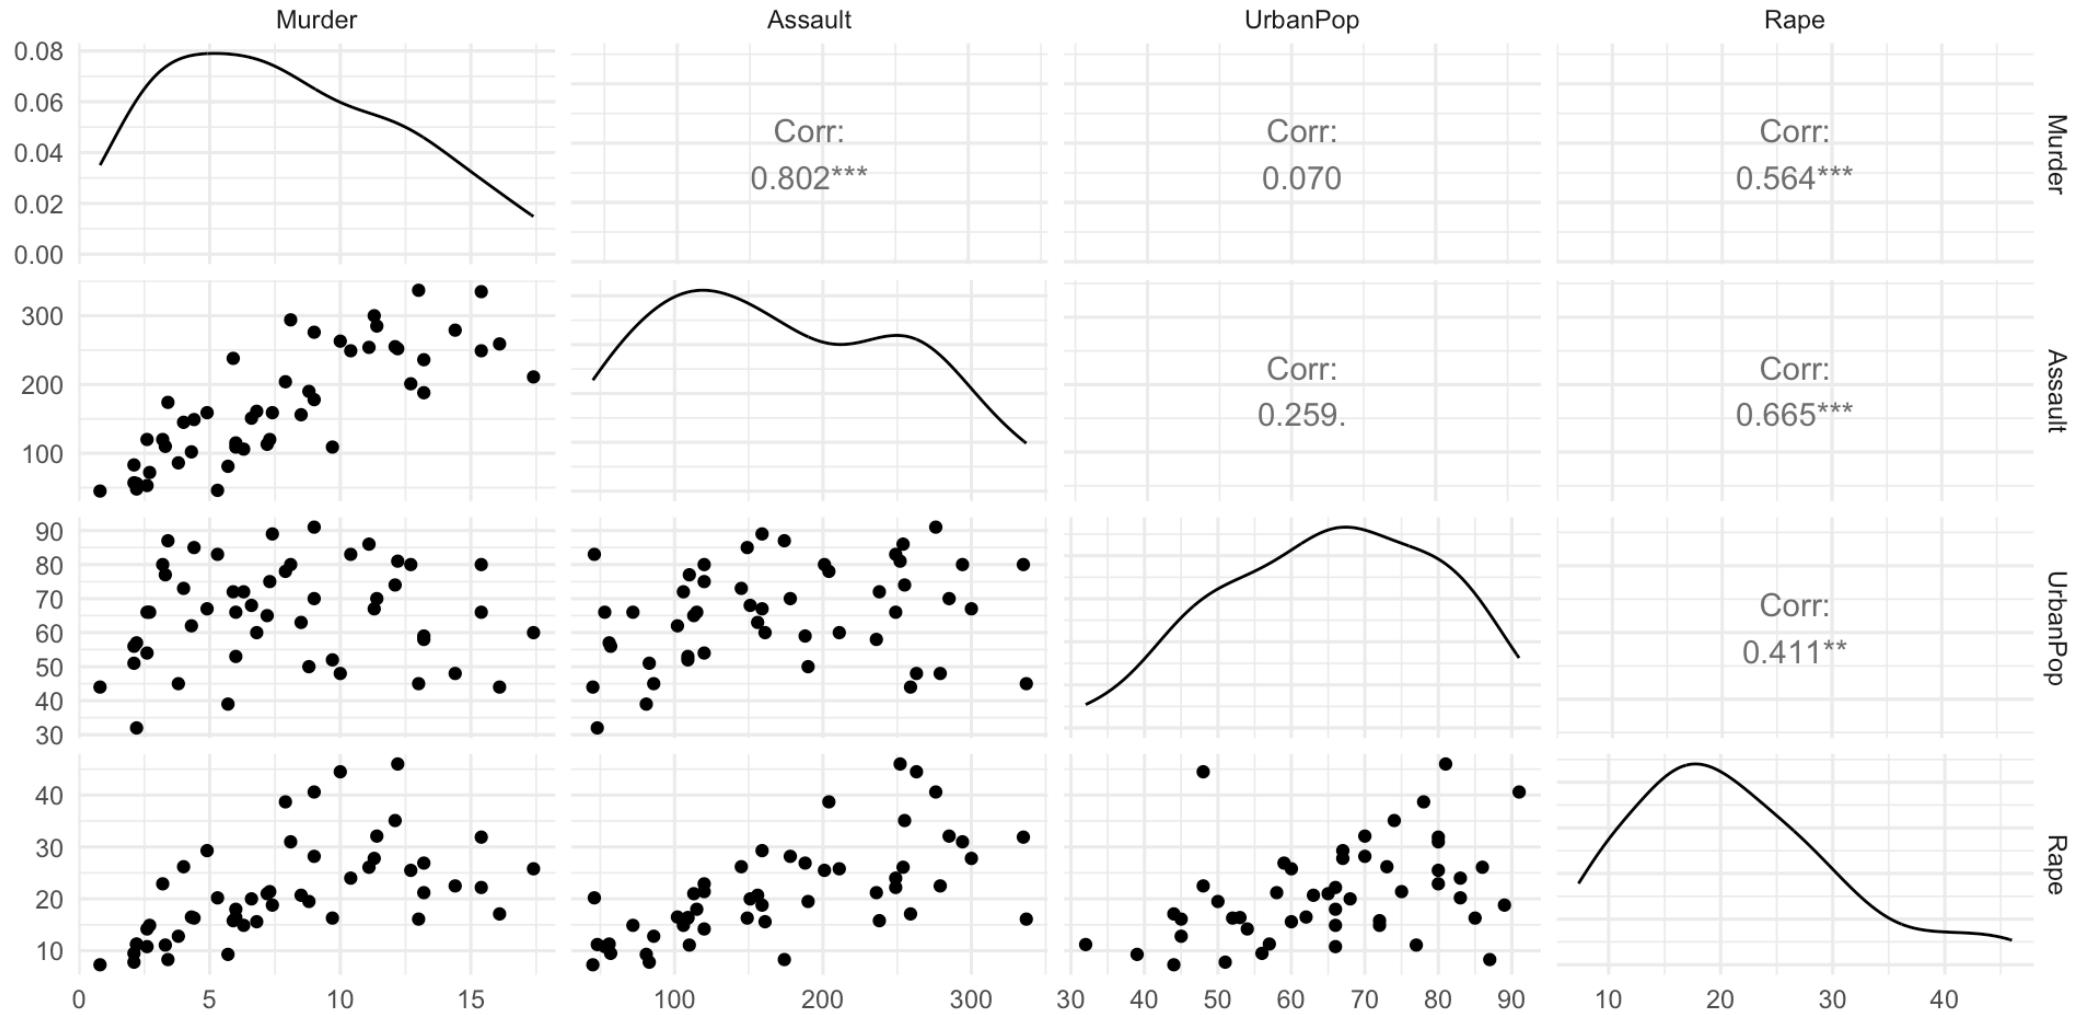


Scatter plot with fitted line and residuals



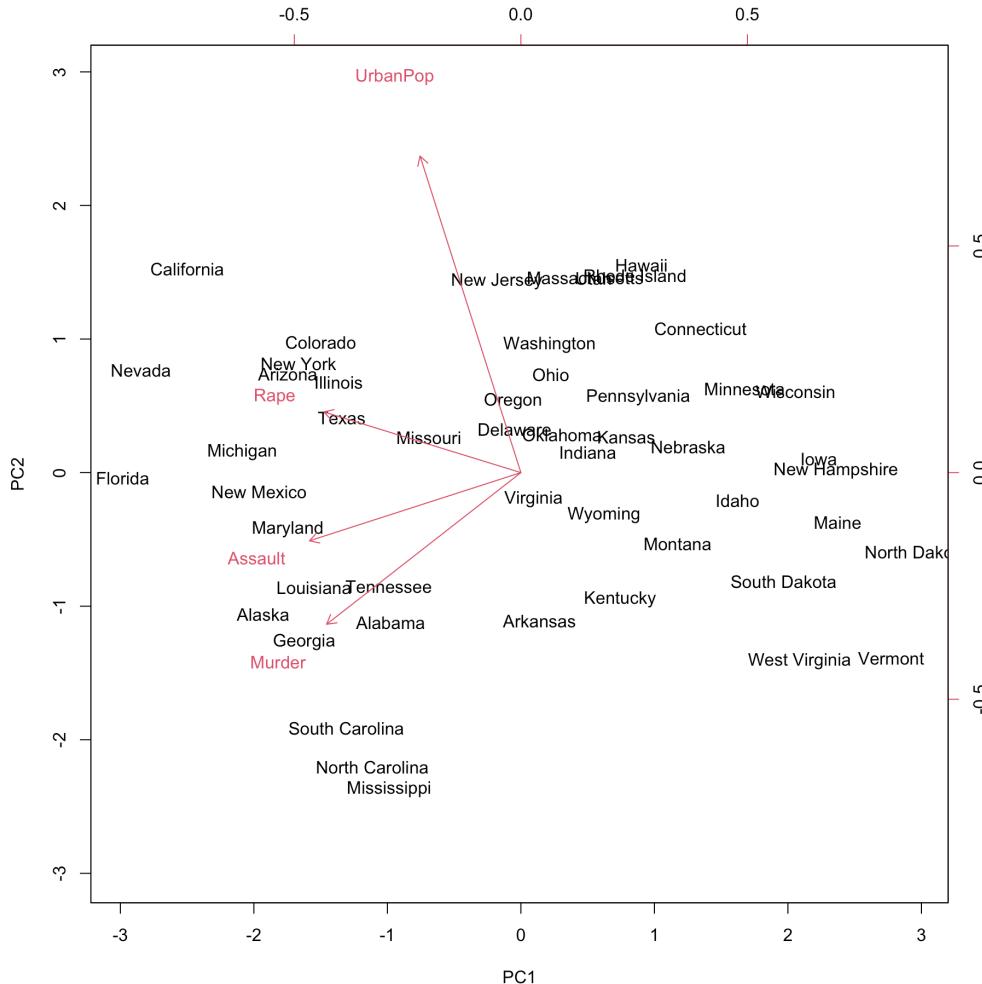
# USArests Example

```
1 library(GGally)  
2 ggpairs(USArests) + theme_minimal()
```



# Biplot of the USArrests

## ► Code



# Preprocessing for PCA

- In **PCA**, it's common to **center** variables by subtracting the mean.
- You can also **standardise** them so all variables have a **standard deviation of 1**.

## Why Standardisation Matters

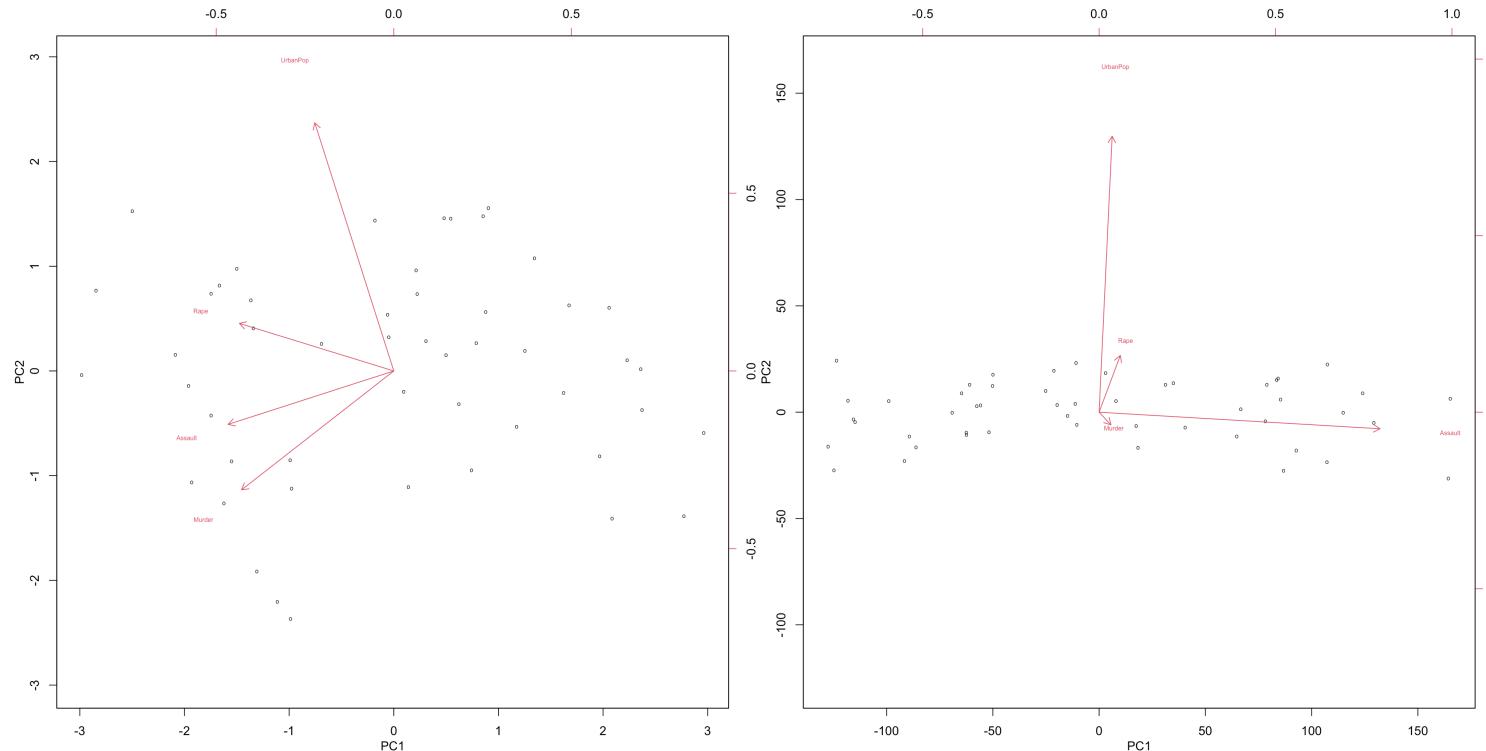
- Variables may have **different units**
- This leads to **different variances**, affecting PCA results.

## Impact on Loadings

- **PCA loadings** give more weight to variables with **higher variance**.
- This can **bias** the analysis if some variables dominate.

## Solution: Standardisation

# Effect of scaling (left) vs unscaled (right) in PCA



- Unscaled

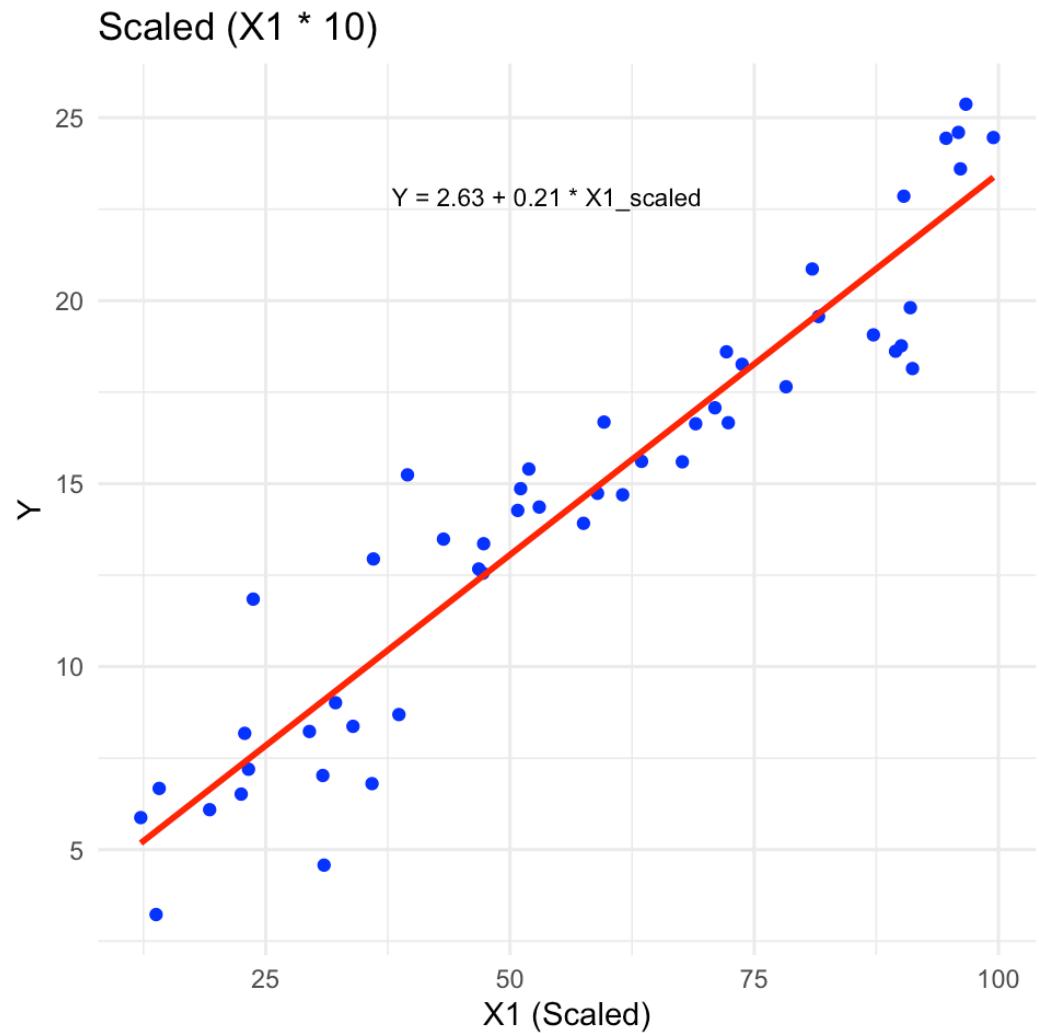
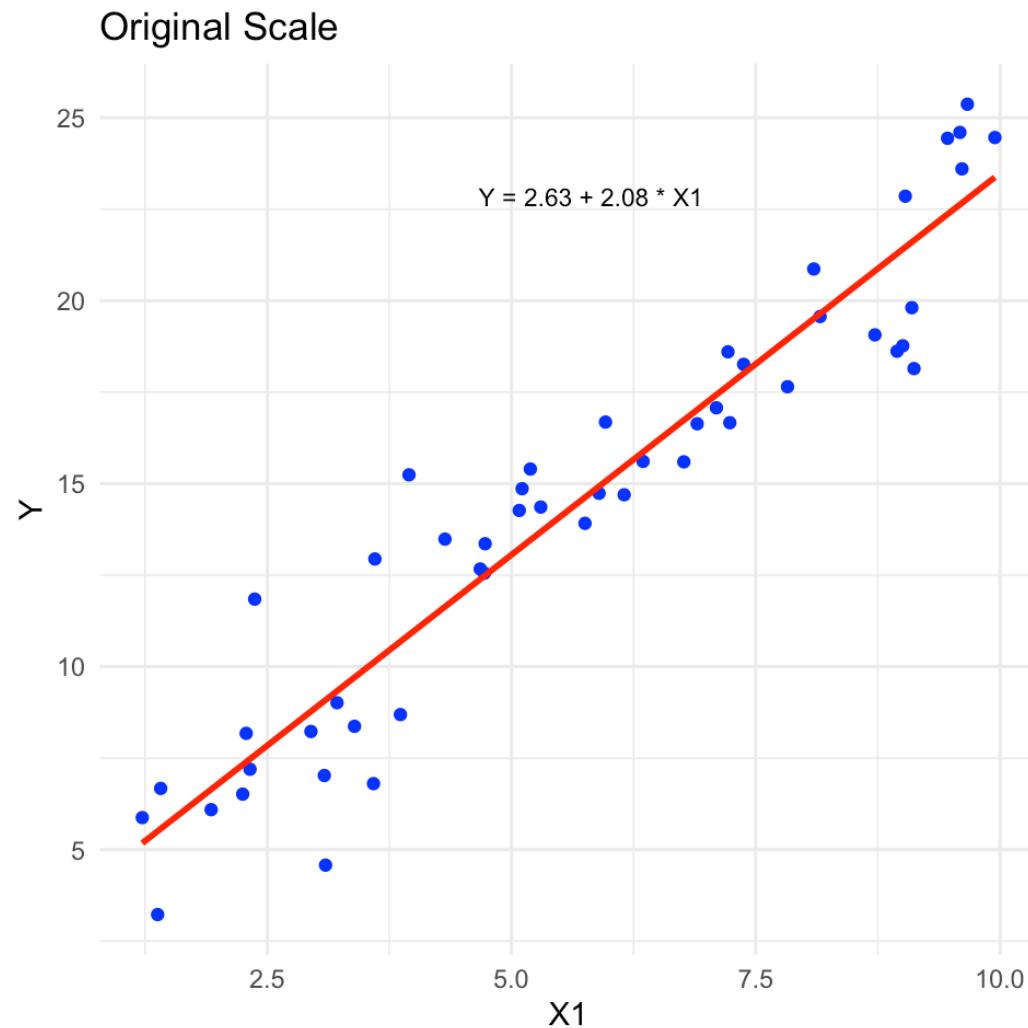


Murder, Rape and assault are reported as the number cases per 100,000 people while UrbanPop is measured in % term.



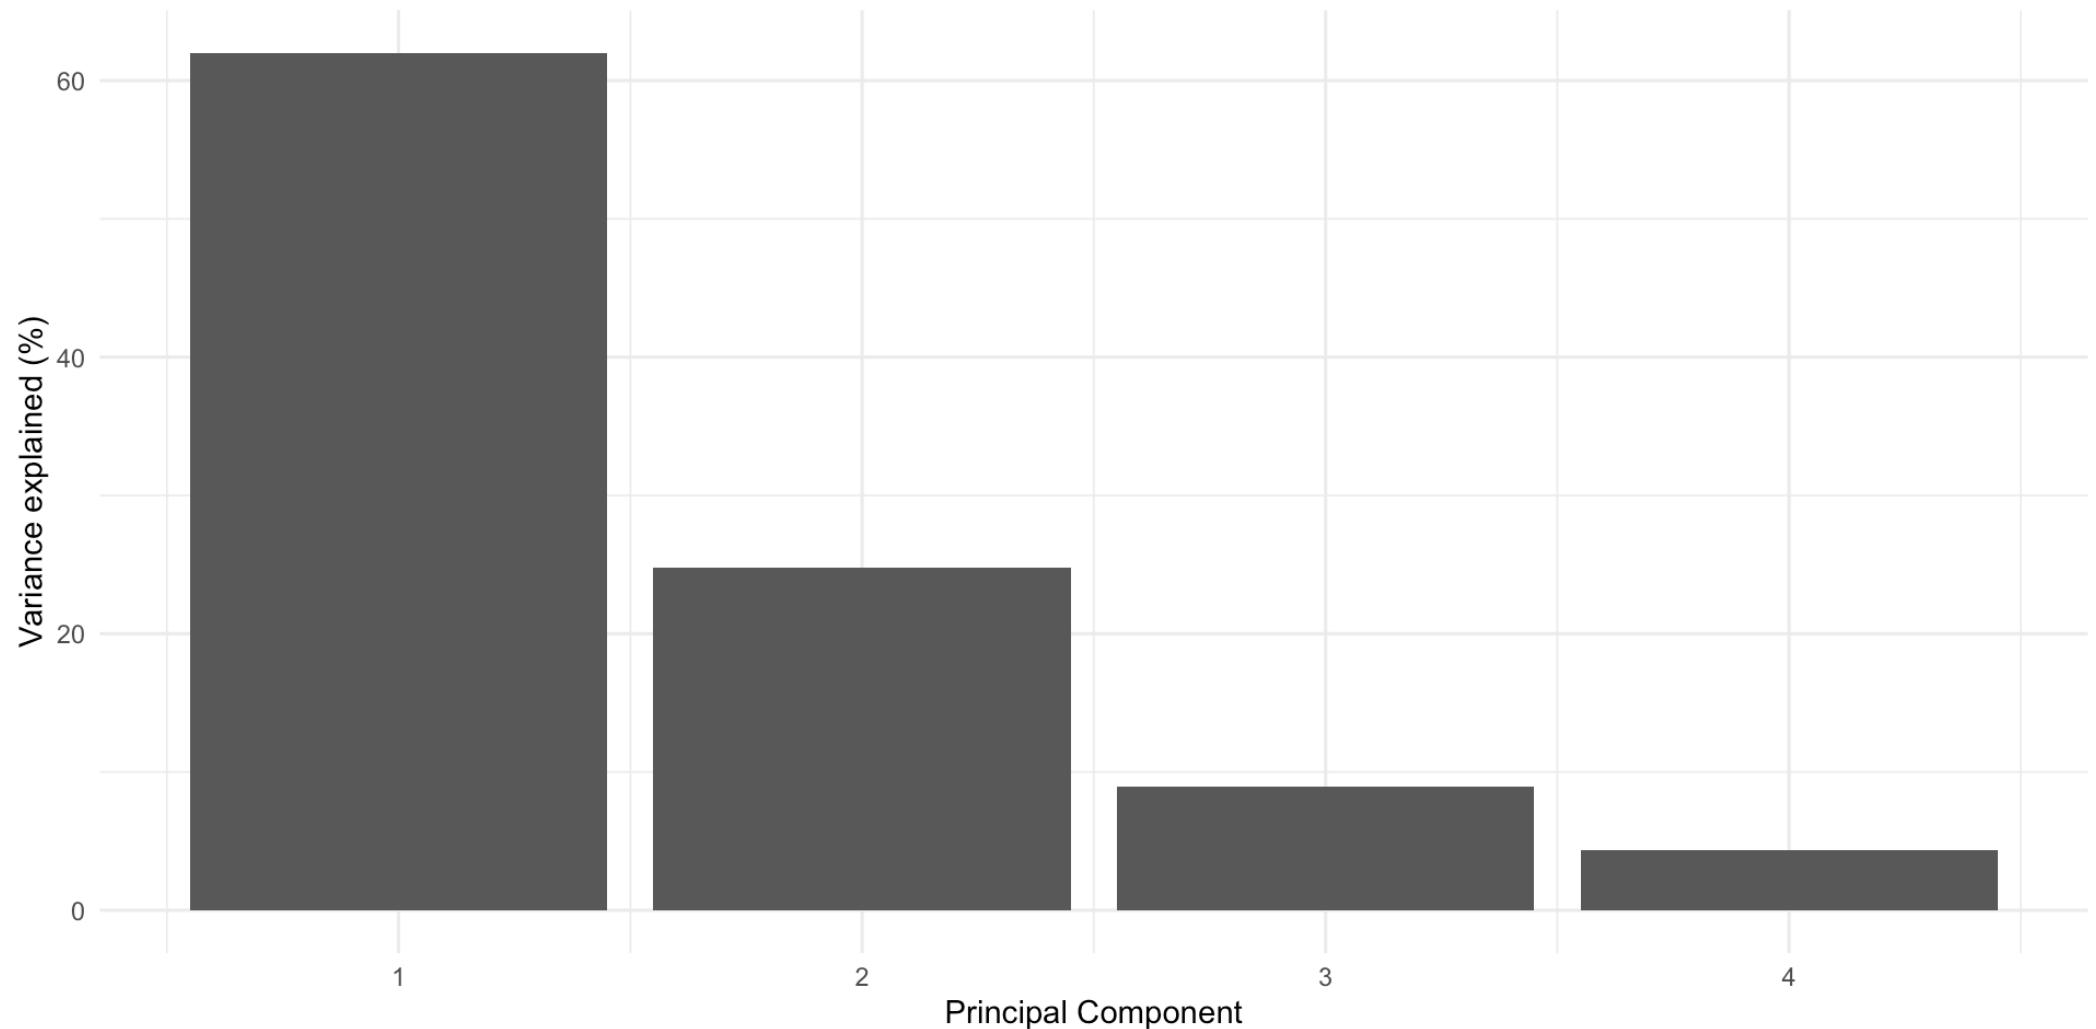
Variances associated with four features are 18.97, 87.73, 6945.16 and 209.5.

# Effect of scaling (left) vs unscaled (right) in linear regression



# Scree plot

► Code



# PCA Properties

- **Unique** and **Global** solution!
- Ordered components
- Best linear dimension reduction possible
- Is not the best for non-linear relationships

## PCA with $K$ -means

- Very common approach to deal with high dimensional data
- Use the first  $M$  principal component scores as inputs into the  $K$ -means algorithm, where  $M \ll p$
- Can help improve the clustering model if the signal in the data can be captured in a few principal components

## PCA with regression

- Use the first  $M$  principal component scores as the predictors in a linear regression model
- We are assuming that a small number of principal components can explain most of the variability in the data as well as the response
- PCA is useful when variables in the data are highly correlated (i.e., collinear)

# Dimension reduction $t$ -SNE

# *t*-SNE

- The name comes from ***t***-distributed **s**tochastic **n**eighbour **e**mbedding
- **Non-linear technique** developed for visualizing high dimensional datasets
- Use local structure in the data to find a low dimensional representation
- Applications include computer security research, music analysis, cancer research, bioinformatics, and biomedical signal processing

# MNIST Example

- 28 by 28 pixel images of handwritten digits
  - 784 features

# Three steps in *t*-SNE

*t*-SNE helps visualize high-dimensional data by preserving local structures in a lower-dimensional space.

## Step 1: Construct High-Dimensional Probability Distribution (Not assessable)

- Define similarities between data points in the high-dimensional space.
- Similar points have a **higher probability** of being neighbors.

Mathematical definition:

$$p_{j|i}(\sigma_i^2) = \frac{\phi(x_j; x_i, \sigma_i^2)}{\sum_{k \neq i} \phi(x_k; x_i, \sigma_i^2)}$$

- $\phi(x; \mu, \sigma)$  denotes the Gaussian density
- Think of  $p_{j|i}(\sigma_i^2)$  as a conditional probability that  $x_i$  would pick  $x_j$  as its neighbour if neighbours were picked in proportion to their probability density under a Gaussian centered at  $x_i$  and with variance  $\sigma_i^2$
- $\sigma_i$  is determined by a parameter called **perplexity**
- The conditional probability can be made symmetric with

$$p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$$

## Step 2: Define Low-Dimensional Probability Distribution (Not assessable)

- Map the data into a lower-dimensional space (e.g., **2D or 3D**).
- Define a new probability distribution to reflect the relationships between data points.
- Use a **Student's t-distribution** (with one degree of freedom) instead of a Gaussian to avoid “crowding” effects.

Mathematical definition:

$$q_{ij} = \frac{(1 + \|y_j - y_i\|_2^2)^{-1}}{\sum_{k \neq i} (1 + \|y_j - y_k\|_2^2)^{-1}}$$

- The above formula comes from the probability density function of **t**-distributions with one degree of freedom (i.e., Cauchy distribution)

## Step 3: Optimizing with Kullback-Leibler (KL) Divergence (Not assessable)

- Adjust the positions of points in the low-dimensional space to **minimize** the difference between high-dimensional and low-dimensional distributions.
- KL divergence is a non-symmetric measure of the difference between two probability distributions
- KL divergence is defined as:

$$\text{KL}(p||q) = \sum_{i \neq j} p_{ij} \log \left( \frac{p_{ij}}{q_{ij}} \right)$$

Think of KL divergence as a measure of how many bits of information is lost when we use  $q$  to approximate  $p$

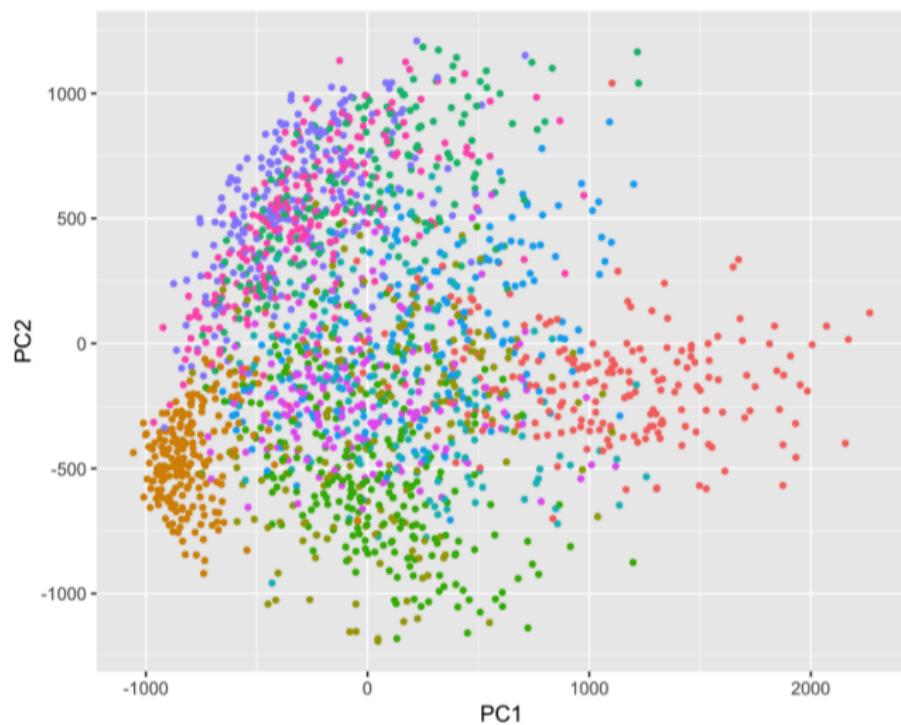
where:

- A lower KL divergence means the lower-dimensional representation **better matches** the original high-dimensional structure.

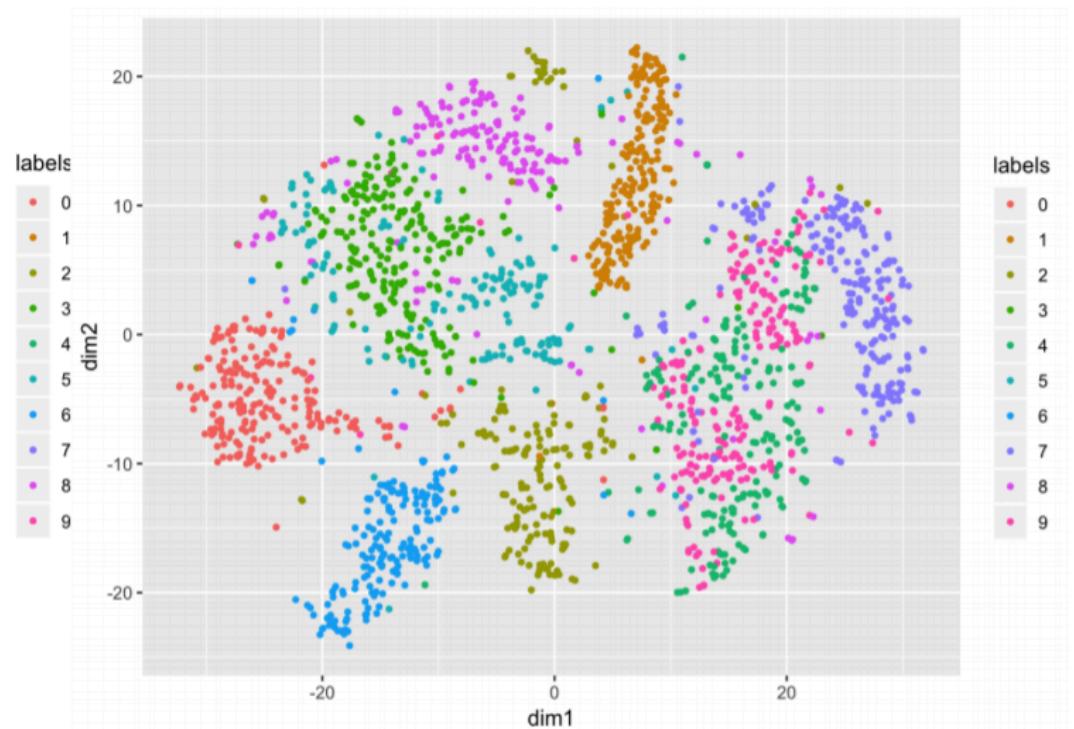
# *t*-SNE vs PCA

MNIST Example

PCA



tSNE



- PCA is a linear method and can only capture linear relationships.
  - PCA is defined by a mathematical formula.
  - The principal components (PCs) from PCA can be interpreted and used for inference.
  - PCA is less computationally intensive
- 
- *t*-SNE is a non-linear technique and is able to find complicated non-linear relationships.
  - *t*-SNE is a probabilistic method – it will give you a different representation every time you run it
  - *t*-SNE is mostly a visualization method. The visualization from *t*-SNE cannot be used for inference.
  - *t*-SNE is more computationally intensive

# Dimension reduction: Multidimensional Scaling (MDS)

# Multidimensional Scaling (MDS)

- Visually represent proximities (similarities or distances) between objects in a lower dimensional space (usually 2 or 3d space)
- The objective of MDS is to take a Matrix of similarities or dissimilarities,  $\mathbf{D}$ , and find projections  $\mathbf{z}_1, \dots, \mathbf{z}_k$  where  $k$  is the desired lower dimension
- The distances are near preserved by optimizing a stress function
- Full data not required

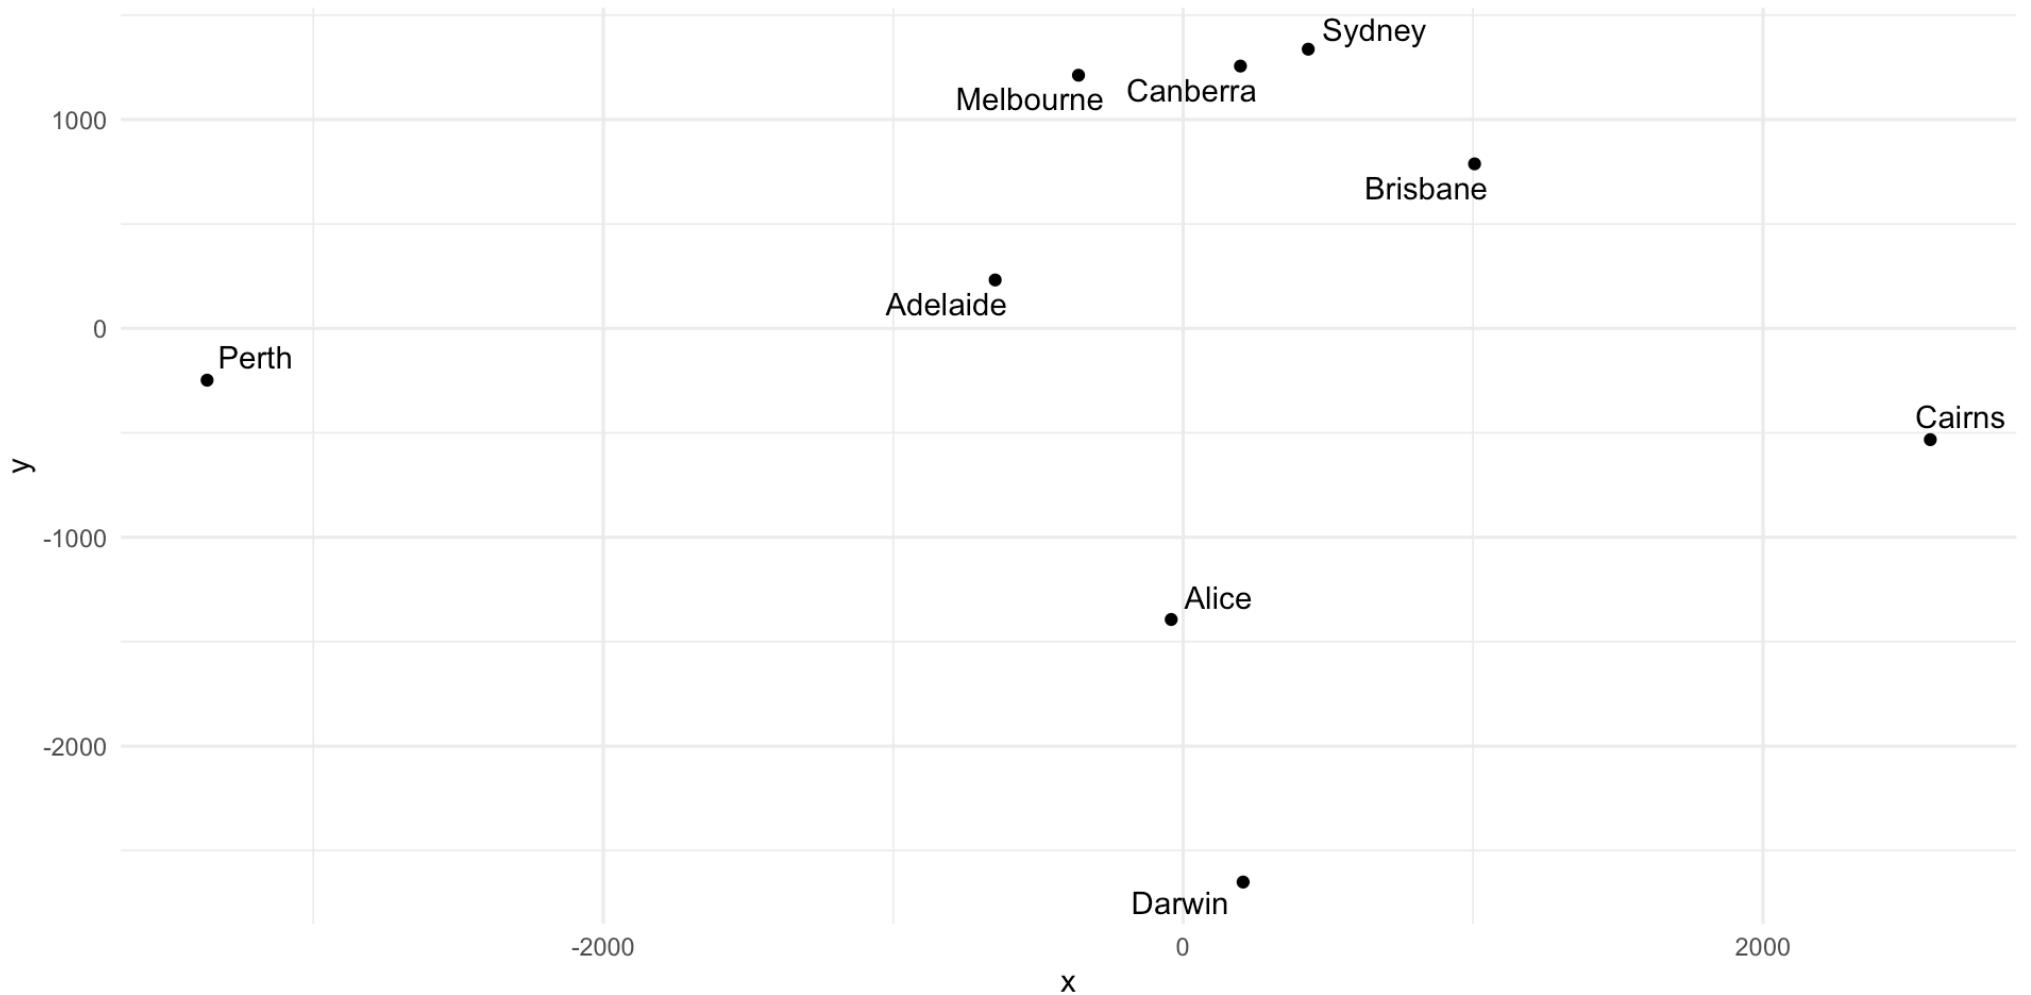
# Multidimensional Scaling (MDS) Example

```
1 cities <- c("Adelaide", "Alice", "Brisbane", "Cairns", "Canberra", "Darwin", "Melbourne", "Perth", "Sydney")
2 city.dist <- matrix(c(0, 1533, 2044, 3143, 1204, 3042, 728, 2725, 1427,
3 1533, 0, 3100, 2500, 2680, 1489, 2270, 3630, 2850,
4 2044, 3100, 0, 1718, 1268, 3415, 1669, 4384, 1010,
5 3143, 2500, 1718, 0, 2922, 3100, 3387, 5954, 2730,
6 1204, 2680, 1268, 2922, 0, 3917, 647, 3911, 288,
7 3042, 1489, 3415, 3100, 3917, 0, 4045, 4250, 3991,
8 728, 2270, 1669, 3387, 647, 4045, 0, 3430, 963,
9 2725, 3630, 4384, 5954, 3911, 4250, 3430, 0, 4110,
10 1427, 2850, 1010, 2730, 288, 3991, 963, 4110, 0),
11 nrow = length(cities), ncol = length(cities))
12 colnames(city.dist) <- rownames(city.dist) <- cities
13 city.dist %>% kbl %>% kable_paper("hover", full_width = TRUE)
```

	Adelaide	Alice	Brisbane	Cairns	Canberra	Darwin	Melbourne	Perth	Sydney
Adelaide	0	1533	2044	3143	1204	3042	728	2725	1427
Alice	1533	0	3100	2500	2680	1489	2270	3630	2850
Brisbane	2044	3100	0	1718	1268	3415	1669	4384	1010
Cairns	3143	2500	1718	0	2922	3100	3387	5954	2730
Canberra	1204	2680	1268	2922	0	3917	647	3911	288
Darwin	3042	1489	3415	3100	3917	0	4045	4250	3991
Melbourne	728	2270	1669	3387	647	4045	0	3430	963
Perth	2725	3630	4384	5954	3911	4250	3430	0	4110

# Multidimensional Scaling (MDS) Example

```
1 mds <- cmdscale(city.dist, k = 2); colnames(mds) <- c("x", "y")
2 mds <- data.frame(mds, City = colnames(city.dist))
3 ggplot(mds, aes(x = x, y = y, label = City)) + geom_point() + ggrepel::geom_text_repel() + theme_minimal()
```



## MDS Stress functions

- These functions attempt to force the lower dimensional projections to preserve the distances in the original data
- Common stress functions

⇒ Least squares  $S_{\text{LS}}(z_1, z_2) = \sqrt{\sum_{i \neq j} (d_{ij} - \|z_i - z_j\|)^2}$

# Comments on MDS

## Benefits & Drawbacks of MDS

- Requires only a distance or dissimilarity matrix, not the full dataset.
- Must choose the number of dimensions  $K$  (similar to the elbow method in a Scree plot).
- Useful for visualizing non-linear data in a lower-dimensional space.

## How to Interpret MDS Maps

- Rotation does not matter – axes and orientation are arbitrary.
- Focus on relative positions – only distances between points are meaningful.
- Objects closer together on the MDS map are more similar based on the input distance matrix.