

STAT5003

Week 7 : Bootstrapping, missing data, and class imbalance

Jaslene Lin

The University of Sydney



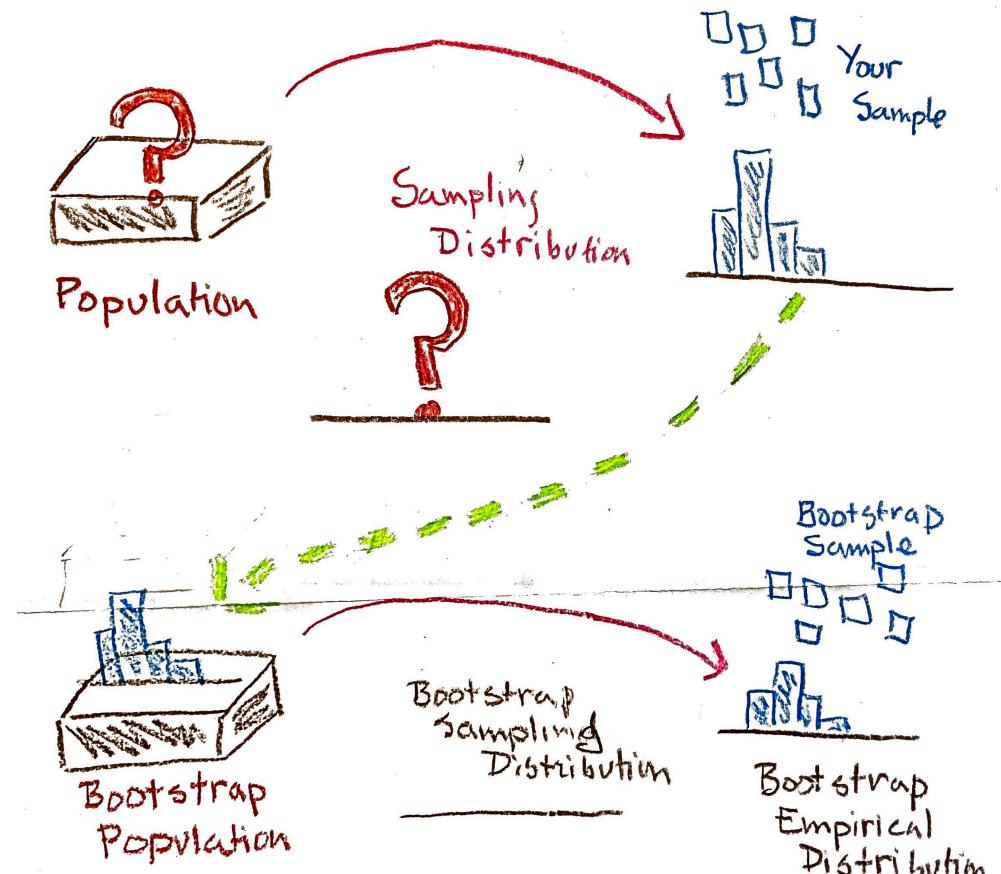
THE UNIVERSITY OF
SYDNEY

Readings and R functions covered

! Important

- **Introduction to Statistical Learning**
 - ➡ Bootstrap covered in Chapter 5
- **R** functions
 - ➡ `naniar:: any_na`
 - ➡ `naniar:: are_na`
 - ➡ `visdat:: vis_dat`
 - ➡ `vidat:: vis_miss`
 - ➡ `mice:: md.pattern`
 - ➡ `mice:: marginplot`
 - ➡ `ROSE:: ovun.sample`
 - ➡ `themis:: step_smote`

Bootstrap



Source

Bootstrap

- The bootstrap is a flexible and powerful statistical tool that can be used to quantify the uncertainty associated with a given estimator or statistical learning method.
- For example, it can provide an estimate of the standard error of a coefficient, or a confidence interval for that coefficient.



Simple example

- Suppose that we wish to invest a fixed sum of money in two financial assets that yield returns of X and Y where X and Y are random quantities
- The goal is to create a portfolio by investing fraction α of our wealth in X and $(1-\alpha)$ in Y
- Want to choose to minimise the total risk of the investment. Mathematically, this involves minimising the variance of our investment portfolio $\text{Var}(\alpha X + (1-\alpha)Y)$
- The solution to this problem (*via calculus*) is

$$\alpha = \frac{\sigma_Y^2 - \sigma_{XY}}{\sigma_X^2 + \sigma_Y^2 - 2\sigma_{XY}} \backslash \text{nonumber} \quad (1)$$

⇒ $\sigma_X^2 = \text{Var}(X)$, $\sigma_Y^2 = \text{Var}(Y)$, and $\sigma_{XY} = \text{Cov}(X, Y)$

A Thought Experiment

- The values of σ_X^2 , σ_Y^2 , and σ_{XY} are unknown but estimates can be computed from the data
- The estimate of α that minimises the variance of the investment can then be computed with

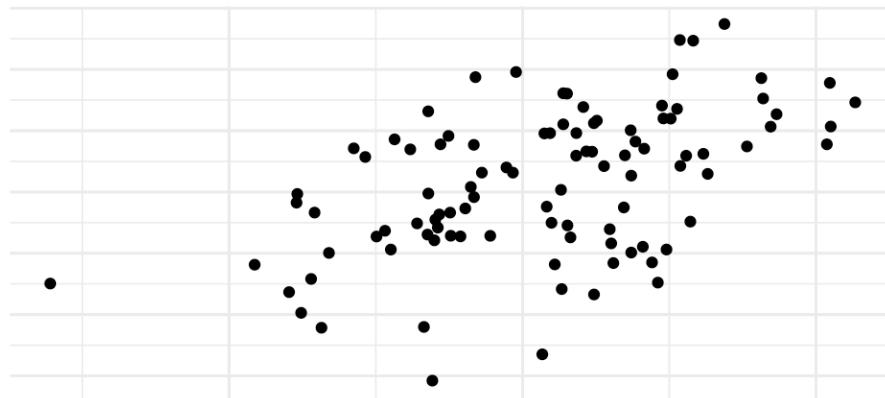
$$\hat{\alpha} = \frac{\hat{\sigma}_Y^2 - \hat{\sigma}_{XY}}{\hat{\sigma}_X^2 + \hat{\sigma}_Y^2 - 2\hat{\sigma}_{XY}} \text{ \textcolor{red}{nonumber}} \quad (2)$$

- Suppose that X and Y can be sampled from the population repeatedly
- You sample **100 pairs of (X, Y)** from the population – this gives you one dataset.
- You calculate an estimate, say $\hat{\alpha}$, from this sample.
- Now, repeat this **sampling process 1000 times**, each time drawing a new random sample of 100 pairs from the population and computing $\hat{\alpha}$.
- You will end up with $\hat{\alpha}_1, \hat{\alpha}_2, \dots, \hat{\alpha}_{1000}$

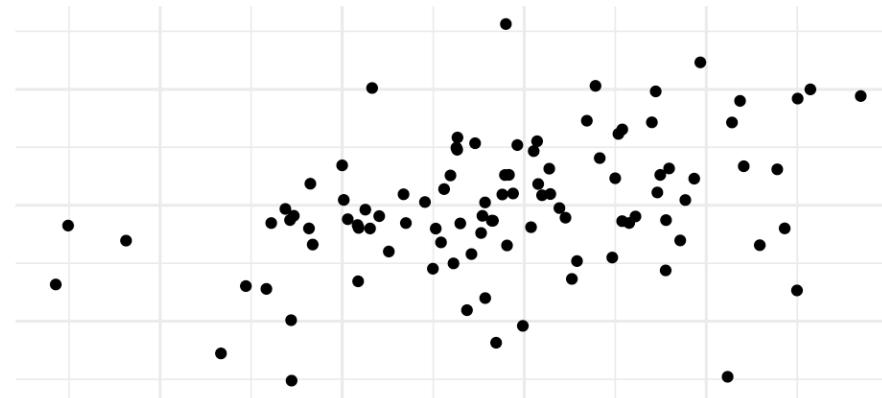
Simulations in R

- Consider a population with $\sigma_X^2 = 1$, $\sigma_Y^2 = 1.5$, and $\sigma_{XY} = 0.5 \Rightarrow \alpha = 2/3$.

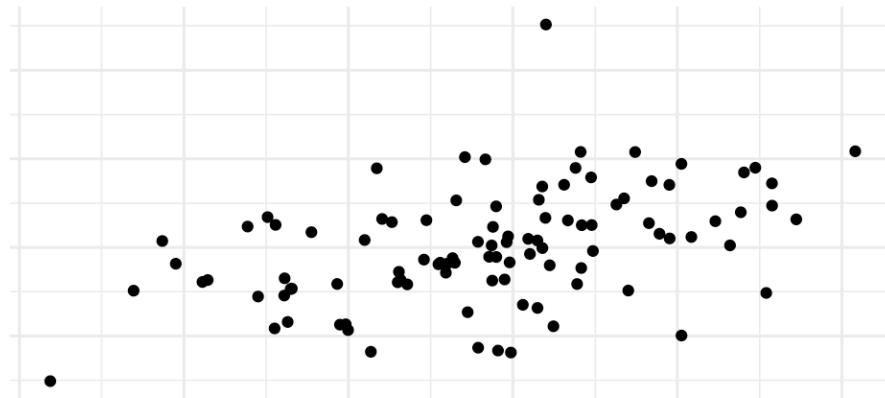
Sample: 1



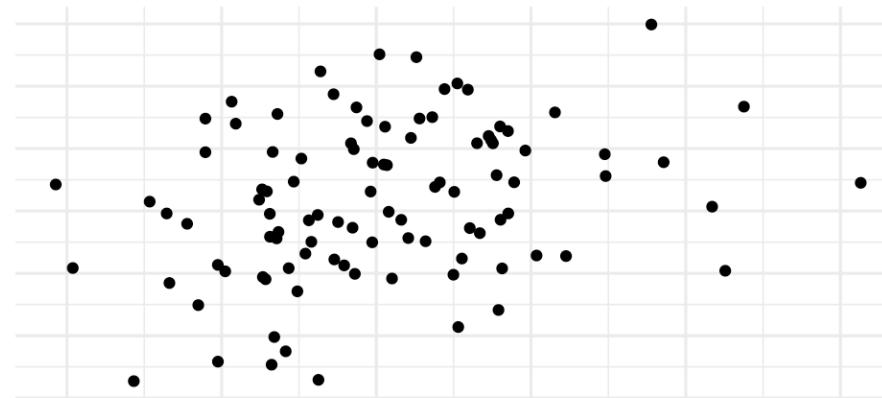
Sample: 2



Sample: 3



Sample: 4



From left to right, top to bottom, the estimates for α are 0.659, 0.683, 0.726, 0.68.

Sampling Variability and the Variance of an Estimator

Note

Key Idea: When we estimate a parameter, the result depends on the sample. If we collecte a different sample, the estimate would likely be different.

- This phenomenon is called **sampling variability**.
- The **variance of an estimator** captures this variability:
 - ➡ If an estimator has **low variance**, repeated samples yield similar results.
 - ➡ If it has **high variance**, the estimate fluctuates a lot between samples.
- Ideally, we want estimators that are:
 - ➡ **Unbiased** (on average, close to the true value)
 - ➡ **Low variance** (stable across samples)
- Estimating the variance of an estimator requires repeatedly sample from the population.

Esmating the Variability of $\hat{\alpha}$

- Consider the mean of all the parameter estimates

$$\bar{\hat{\alpha}} = \frac{1}{1000} \sum_{k=1}^{1000} \hat{\alpha}_k = 0.6662595$$

- This is close to the true value of 0.6666667
- Estimate of the standard error using the standard deviation of all the estimates

$$\sqrt{\frac{1}{1000 - 1} \sum_{k=1}^{1000} (\hat{\alpha}_k - \bar{\hat{\alpha}})^2} = 0.0760217$$

- This gives an intuitive description of the reliability of the estimator
 - For a random sample the estimate would vary around the true value by 0.0760217

Application in Reality

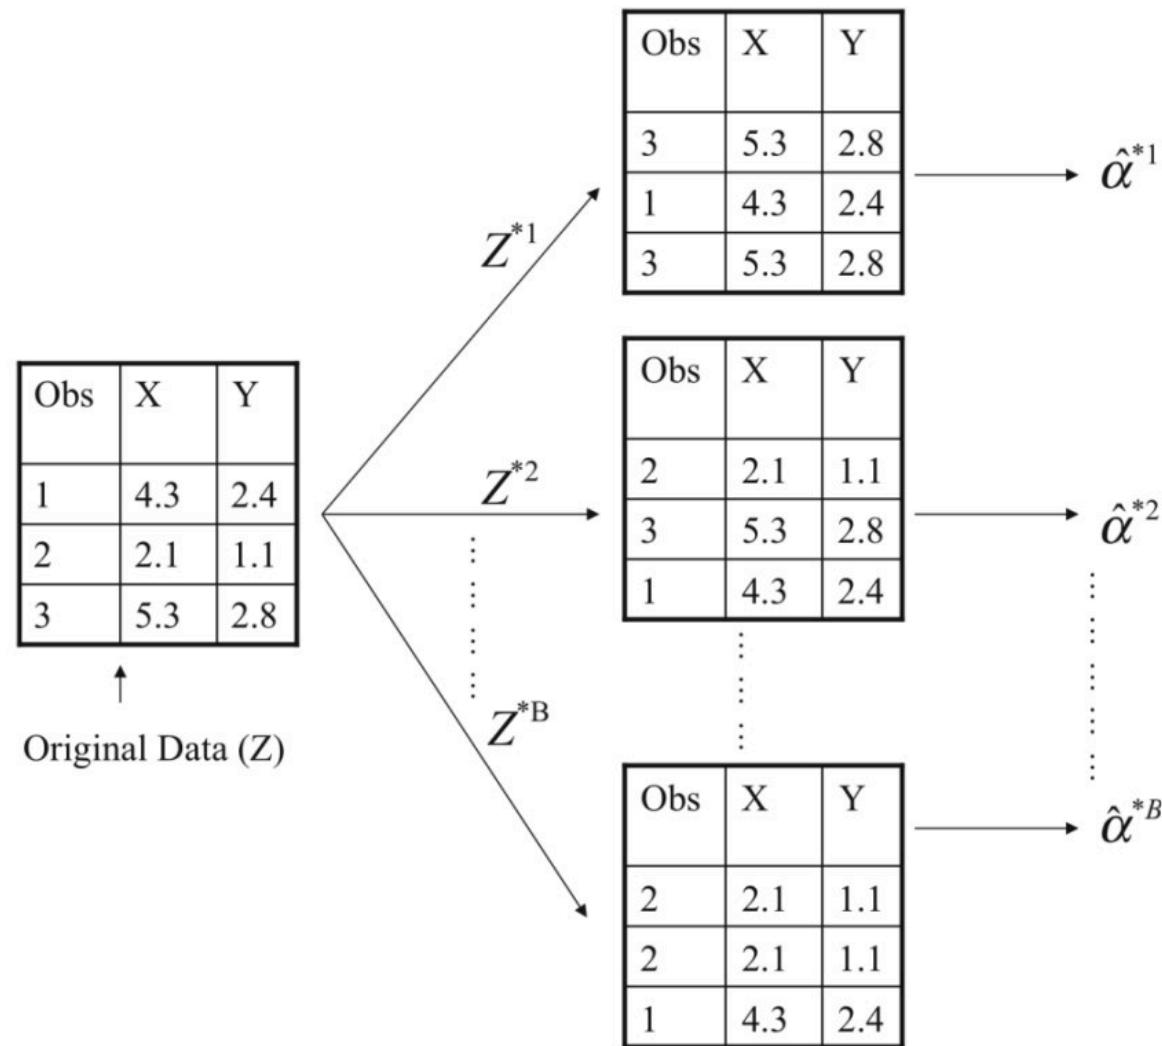
- We can't actually sample repeatedly from the population
 - ➡ In reality, we **only have one dataset**, not the entire population to draw from.

Bootstrap attempts to mimic this process

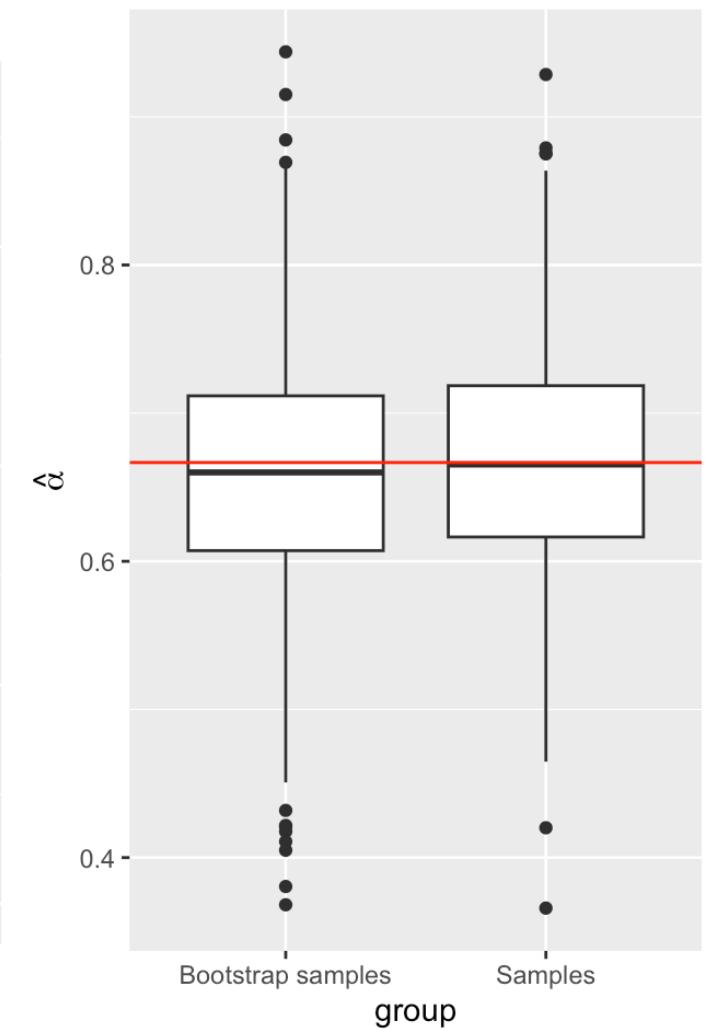
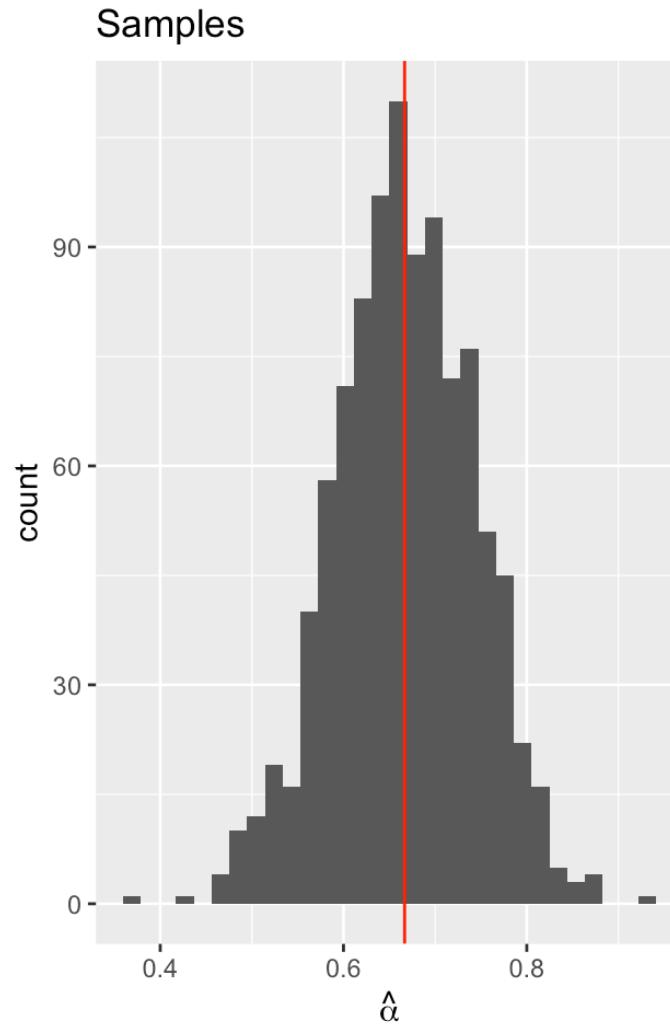
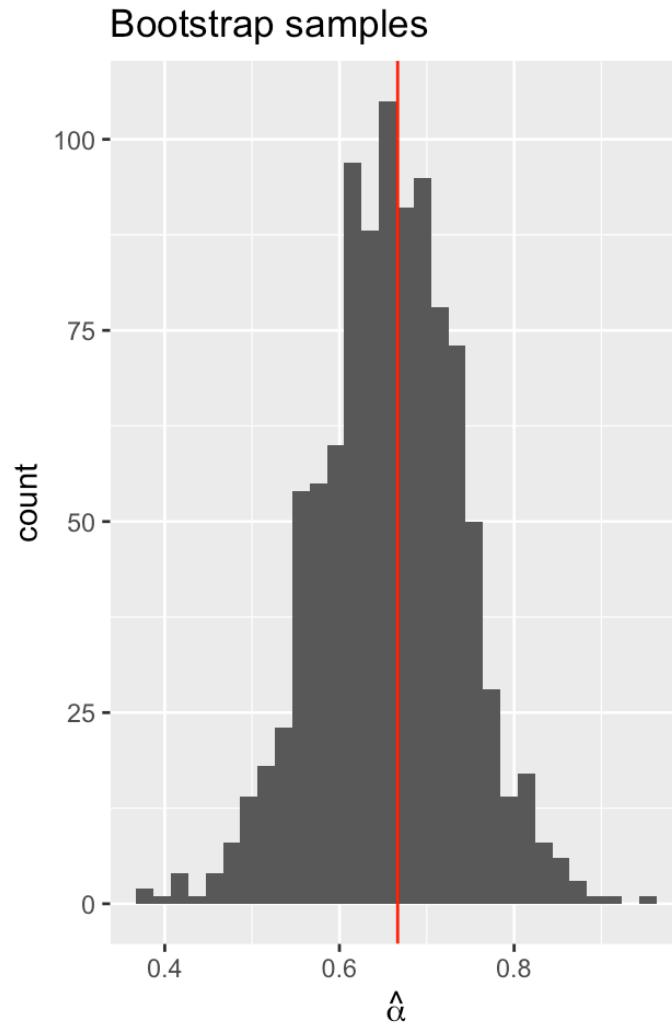
- Instead of sampling new independent observations from the population
 - ➡ Re-sample observations from the data *with replacement*
 - ➡ Some observations may appear more than once while others not at all

Bootstrap resampling algorithm

- Essentially sampling with replacement from the original sample



Results bootstrap vs population



Missing data

The best thing to do with missing data is to not have any.

--- Gertrude Mary Cox (statistician)

- Working with real-world data = working with missing data

Any Missing Data?

```
1 library(naniar)
2
3 glimpse(airquality)
```

```
Rows: 153
Columns: 6
$ Ozone    <int> 41, 36, 12, 18, NA, 28, 23, 19, 8, NA, 7, 16, 11, 14, 18, 14, ...
$ Solar.R  <int> 190, 118, 149, 313, NA, NA, 299, 99, 19, 194, NA, 256, 290, 27...
$ Wind     <dbl> 7.4, 8.0, 12.6, 11.5, 14.3, 14.9, 8.6, 13.8, 20.1, 8.6, 6.9, 9...
$ Temp     <int> 67, 72, 74, 62, 56, 66, 65, 59, 61, 69, 74, 69, 66, 68, 58, 64...
$ Month    <int> 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, ...
$ Day      <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, ...
```

```
1 # str(airquality)
2
3 any_na(airquality)
```

```
[1] TRUE
```

```
1 head(are_na(airquality$Solar.R))
```

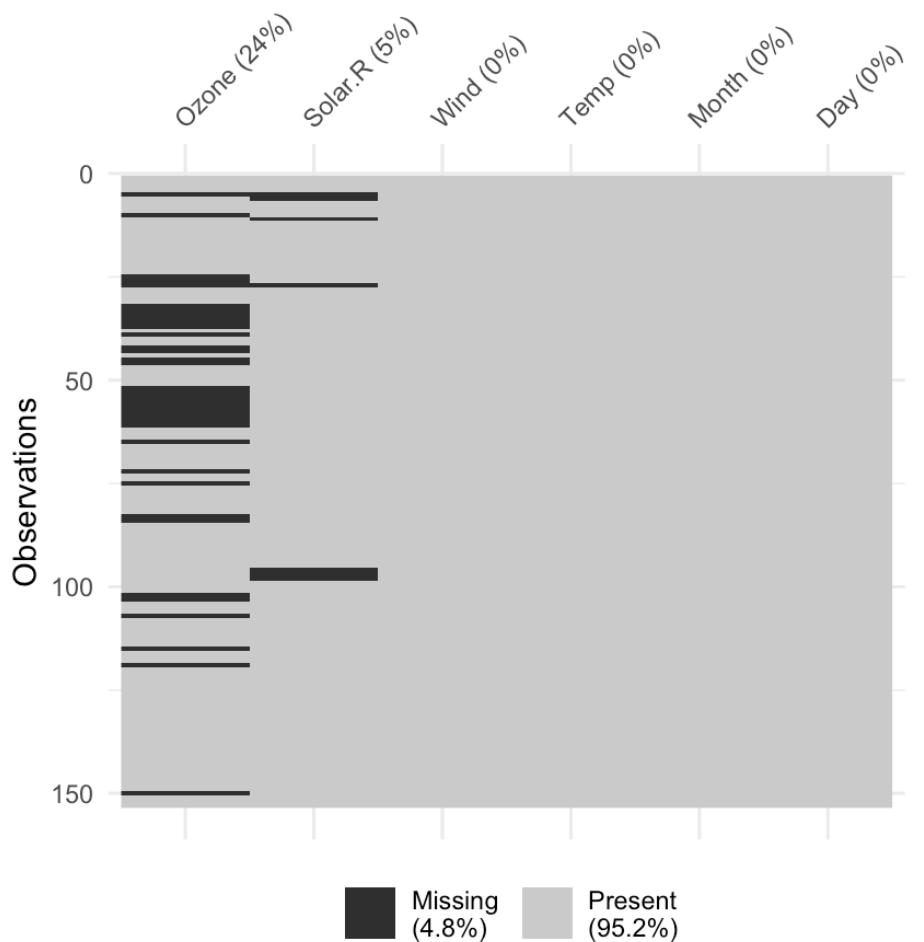
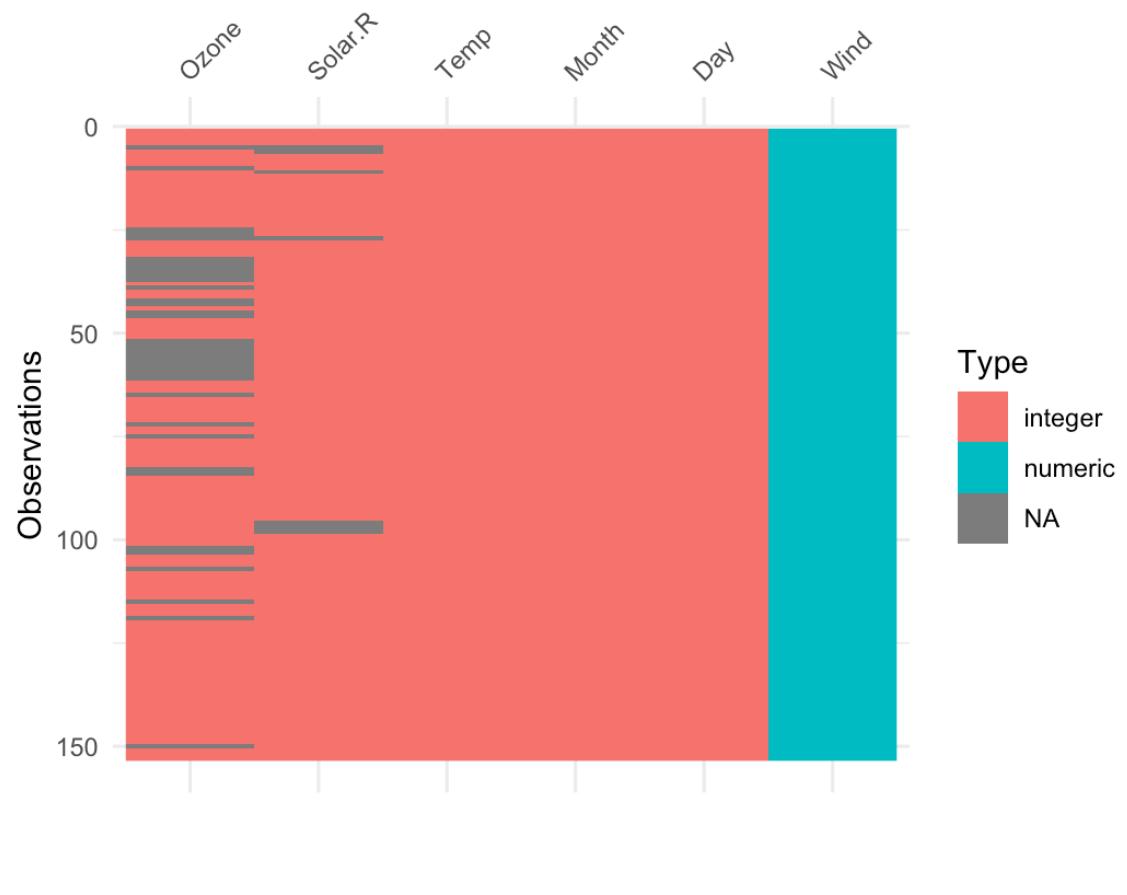
```
[1] FALSE FALSE FALSE FALSE  TRUE  TRUE
```

```
1 any_na(iris)
```

```
[1] FALSE
```

Visualising Missingness

► Code



Missing Mechanisms

1. Missing Completely at Random (MCAR)

- Missingness has no association with any data you have **observed**, or **not observed**. Missingness occurs **purely by chance**.
- Example:

Test	Special Consideration
65	FALSE
NA	TRUE
80	FALSE
70	FALSE
NA	FALSE

- Handling
 - ➡ Imputation is advisable
 - ➡ Deleting observations will not bias the results, but may reduce the model power due to the reduced sample size thus limiting inference.

Missing Mechanisms

2. Missing at Random (MAR)

- Missingness depends on **data observed**, but **not data unobserved**.
- **Example:**

Test	Special Consideration	Health
65	FALSE	Good
NA	TRUE	Poor
80	FALSE	Good
70	FALSE	Good
NA	FALSE	Poor

- **Handling:**
 - ➡ Use **imputation** to adjust for missing data.

Mechanism for missing data

3. Missing Not at Random (MNAR)

- Missingness is related to unobserved data
- **Example:**

Test	Special Consideration	Health
65	FALSE	Good
NA	TRUE	NA
80	FALSE	Good
70	FALSE	Good
NA	FALSE	NA

- **Handling**

➡ Requires **modeling the missingness mechanism** to correct biases.

Identifying different types of missingness

- It's often **difficult to determine the missingness mechanism** using statistical tests alone.
- Instead, rely on:
 - ➡ **Domain knowledge** and understanding of the **data collection process**
 - ➡ **Predictive models**: Fit a classification model to predict missingness—if successful, data may be MAR.
- For testing **MCAR**:
 - ➡ Use tests like `mcar_test()` from the `naniar` package.
 - ➡ Based on Little's MCAR test (Little, 1988), which assumes multivariate normality.

Dealing with missing values

- For categorical data, “missing” can be a category
 - ➡ For example, in a survey poll, if someone does not want to disclose who they want to vote for, it then can be in the category “undecided”
- Delete data with missing value. Two options:
 - ➡ Omit the variable with missing data
 - ➡ Omit the observation with missing data
 - ➡ Drawbacks are that you might be throwing away valuable information, or inadvertently introduce bias into the data
- **Impute: fill in the missingness**

Single imputation

- Single imputation replaces the missing value with a single value
- Examples:
 - ➡ Replace the missing values of a feature with the mean/median value of that feature
 - ➡ Use a predictive method for filling in the missing values, e.g., regression trees, kNN
 - ➡ Replace the missing value with the last observed value for that feature
- **Drawbacks:**

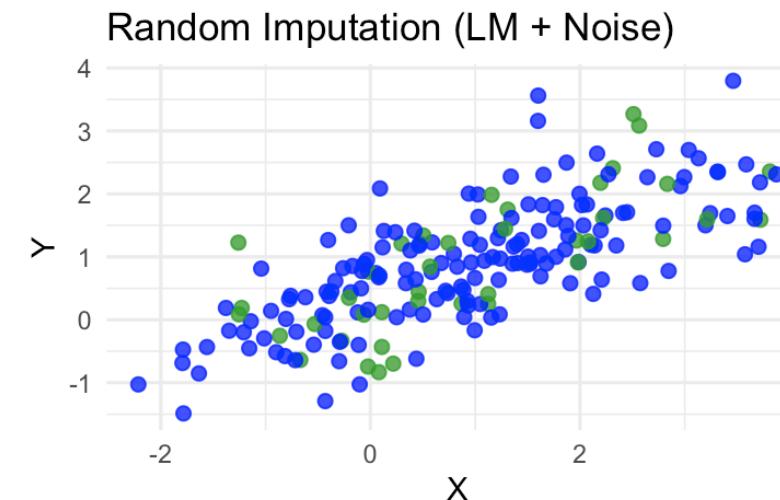
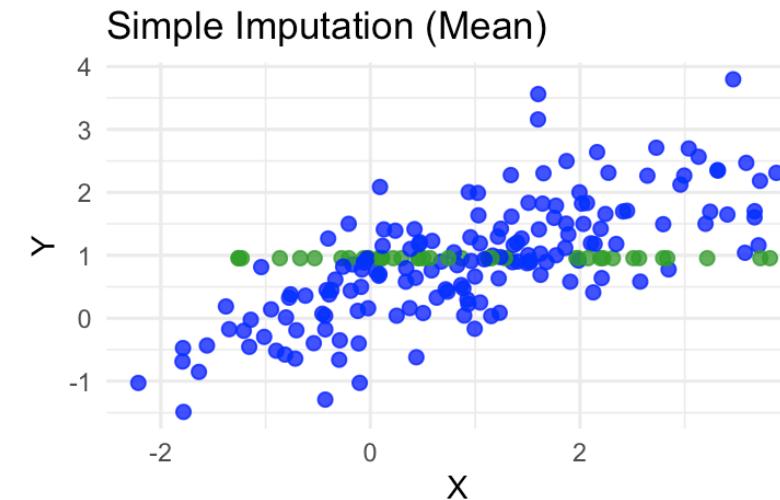
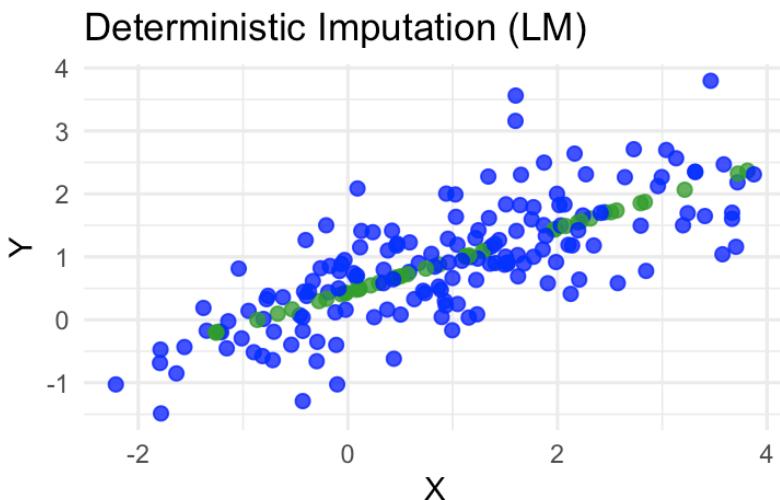
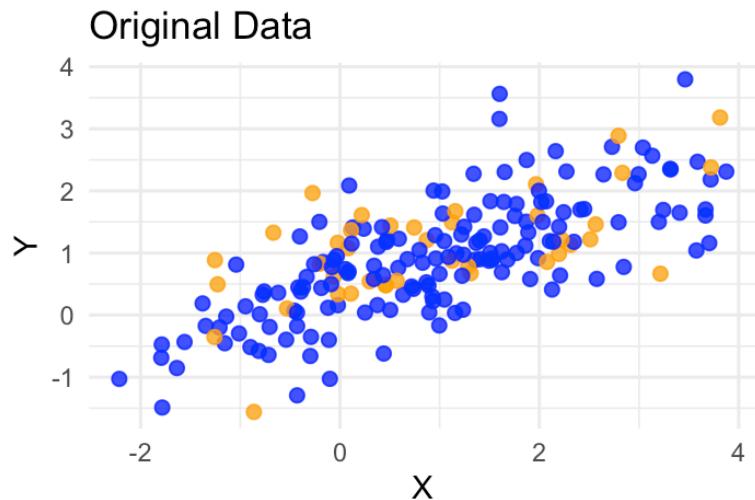
With single imputation, once the missing data are added back, it is treated as valid observed data, hence the uncertainty in the missing value data is lost

Multiple imputation

- Multiple imputation accounts for uncertainty in the imputation process
- Generally follows three steps:
 - ➡ Impute the data k times (this can be done using a single imputation method)
 - ➡ Perform analysis (e.g., regression/classification) on each of the k imputed data sets
 - ➡ Pool the k results together
- Multiple Imputation by Chained Equation (MICE) is a popular method
 - ➡ See [mice: Multivariate Imputation by Chained Equations in R](#)

Basic impute, deterministic imputation, random imputatation

► Code



Other practical suggestions

- It is highly recommended that you visualize your data to look for patterns of missingness
- Be wary of variables with high proportion of missing data. However, this might not be a problem if imputation is applicable and performs well
- Some algorithms can cope with missingness (e.g., decision trees) and so you may not need to do imputation
- If you believe the pattern of missingness is informative, you can include it as a dummy variable

R packages for dealing with missing values

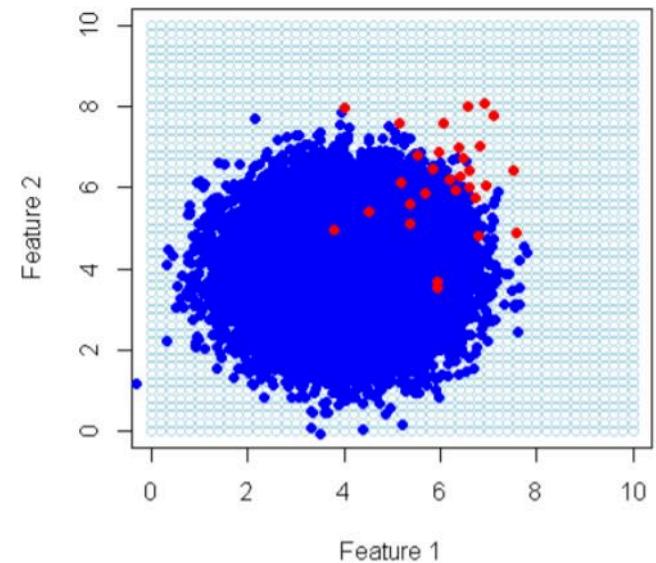
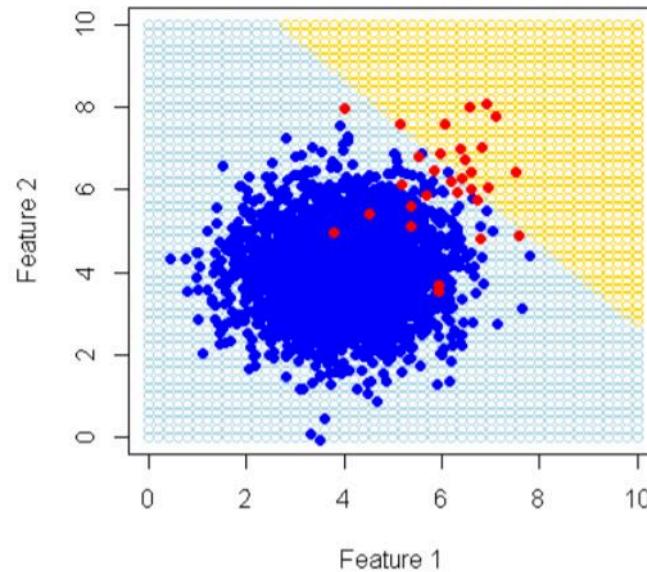
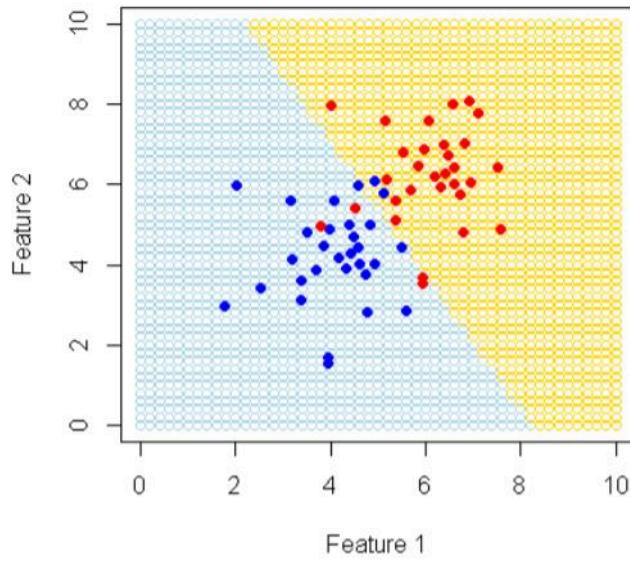
- Impute (Bioconductor package)
 - ➡ KNN imputation, written for microarray data
- MICE
 - ➡ Multiple imputation
- missForest
 - ➡ Uses Random Forest (in upcoming tree-based module) to predict the missing values
 - ➡ Can be used for continuous and categorical data
- Amelia II
 - ➡ Multiple imputation

Class imbalance

Class imbalance

- Let's say we have a classification problem to detect credit card fraud, but only 1% of transactions are fraud.
- If you use accuracy as the metric to optimise, then just classifying every transaction as not-fraud will get you to 99% accuracy!

Inspect a SVM



- Notice the negative (blue) class swamps the classifier
- Boundary moves and disappears favouring only the negative class

Handelling Class Imbalance

1. Use different metrics

- F_1 score

$$F_1 = \frac{2\text{TP}}{2\text{TP} + \text{FP} + \text{FN}} = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

- Area Under Curve (AUC) for the ROC
- Cohen's Kappa

What is Cohen's Kappa metric?

$$\kappa = \frac{P_o - P_e}{1 - P_e}$$

Where:

- P_o = Observed accuracy
- P_e = Expected accuracy generated simply by chance

Interpretation Guide

Kappa Value	Strength of Agreement
< 0	Less than chance
0.01–0.20	Slight
0.21–0.40	Fair
0.41–0.60	Moderate
0.61–0.80	Substantial
0.81–1.00	Almost perfect

Example: Binary Classification – Spam Detection

	Actual: Spam	Actual: Not Spam
Predicted: Spam	30	15
Predicted: Not Spam	10	45

- Observed accuracy:

$$P_o = \frac{30 + 45}{100} = 0.75$$

- Class proportions:

➡ Predicted Spam: (45%); Predicted Not Spam: (55%)

➡ Actual Spam: (40%); Actual Not Spam: (60%)

- Expected accuracy:

$$P_e = (0.4 \times 0.45) + (0.6 \times 0.55) = 0.18 + 0.33 = 0.51$$

- Kappa:

$$\kappa = \frac{0.75 - 0.51}{1 - 0.51} = \frac{0.24}{0.49} \approx 0.49$$

Cohen's Kappa for Class Imbalance Classification

- Example

- ➡ 100 samples with 2% positive class and 98% negative class
- ➡ A classifier classifies all samples as negative class
- ➡ Lead to accuracy 98%, but what about other performance measures?
- ➡ Sensitivity: 0; Specificity: 1; Cohen's kappa: 0

You can compute Kappa in R using `caret::confusionMatrix`

```
Confusion Matrix and Statistics

Reference
Prediction  positive negative
positive      0        0
negative      2       98

Accuracy : 0.98
95% CI  : (0.9296, 0.9976)
No Information Rate : 0.98
P-Value [Acc > NIR] : 0.6767

Kappa : 0

McNemar's Test P-Value : 0.4795

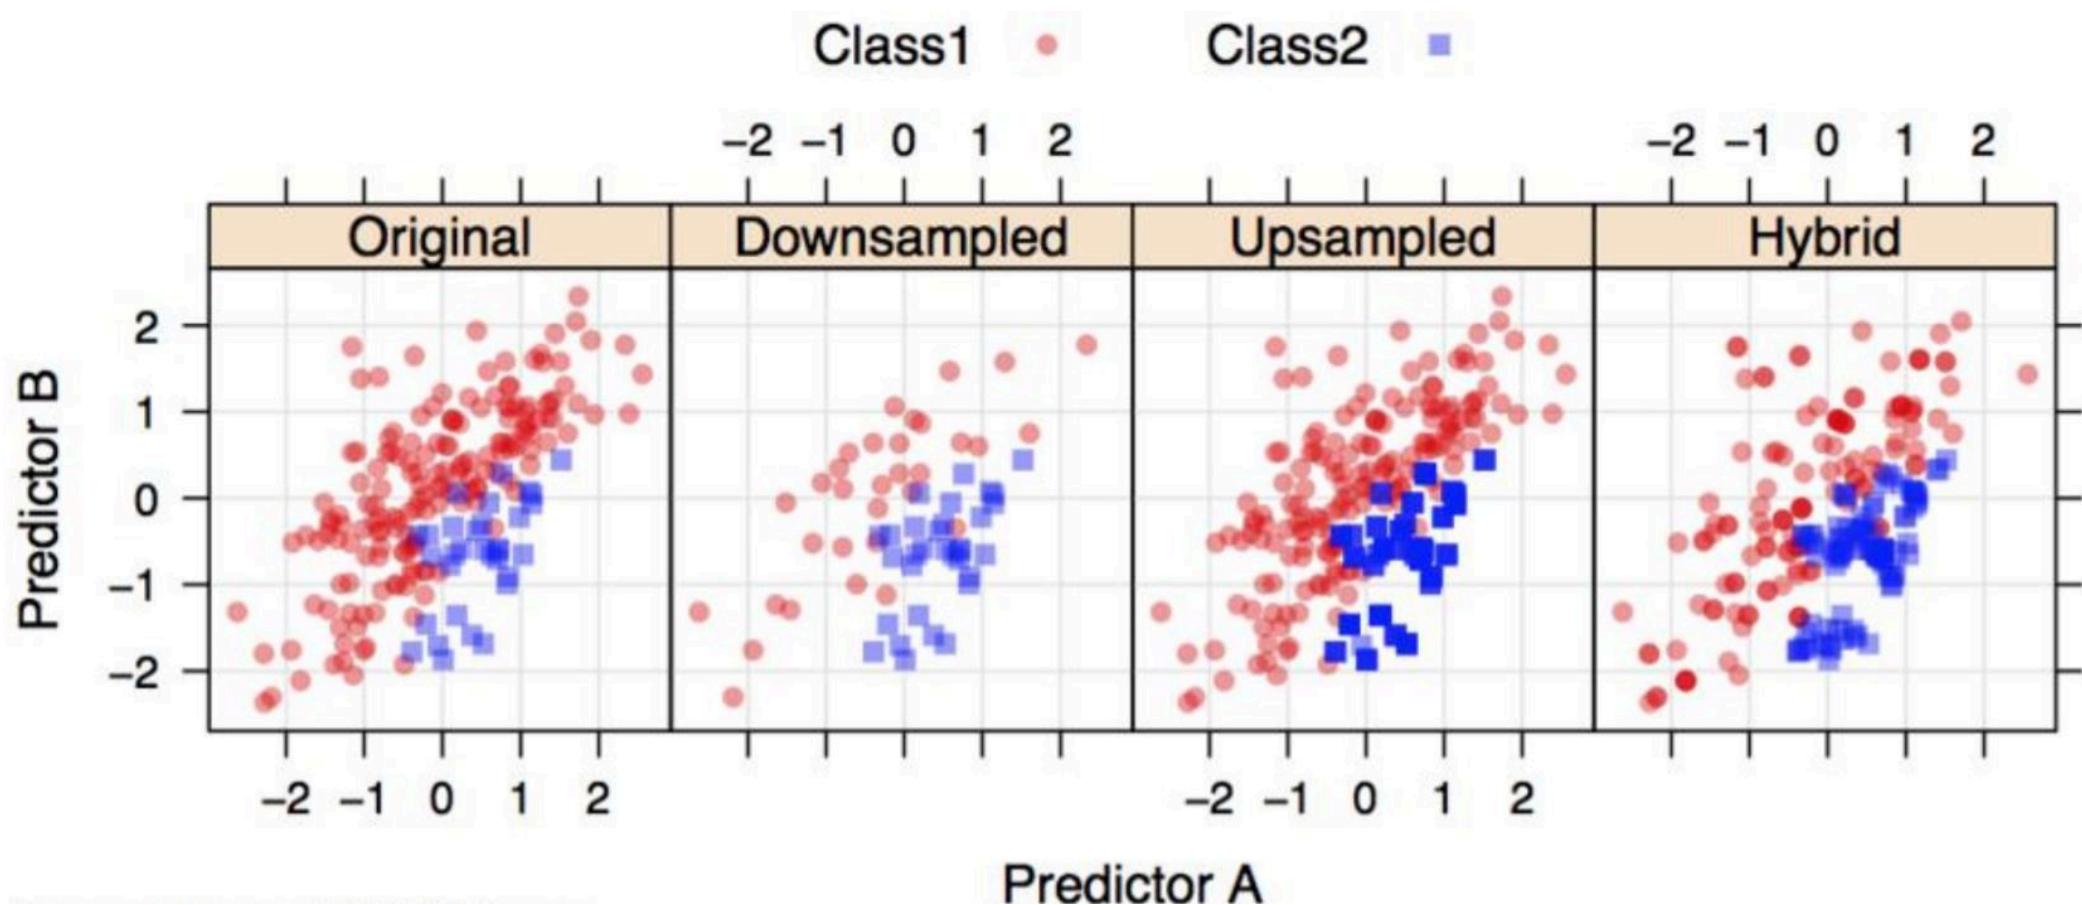
Sensitivity : 0.00
Specificity : 1.00
Pos Pred Value : NaN
Neg Pred Value : 0.98
```

Prevalence : 0.02
Detection Rate : 0.00

Kappa
0

Handling Class Imbalance

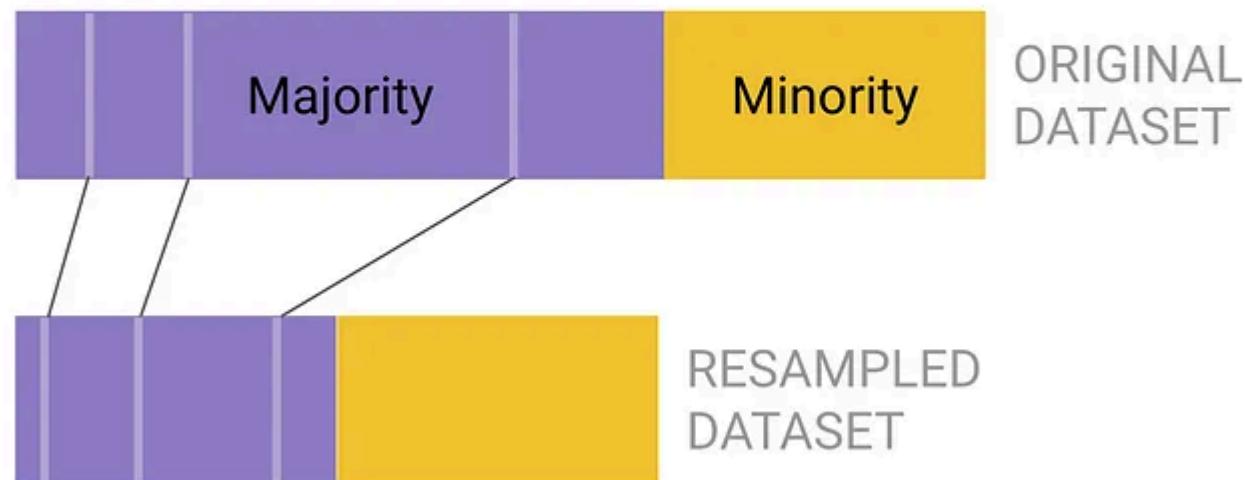
2. Modify the Data



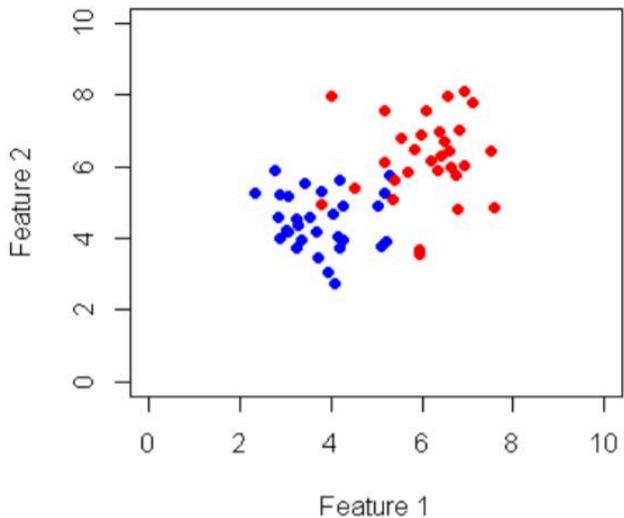
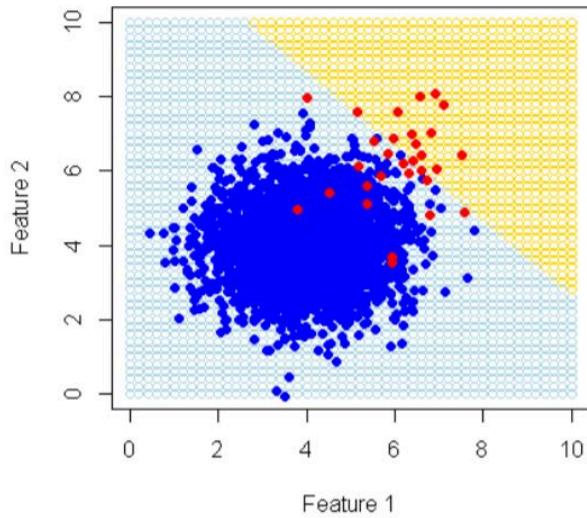
Down-sampling

Undersample majority class

TRAVIS TANG



Travis Tang from Medium

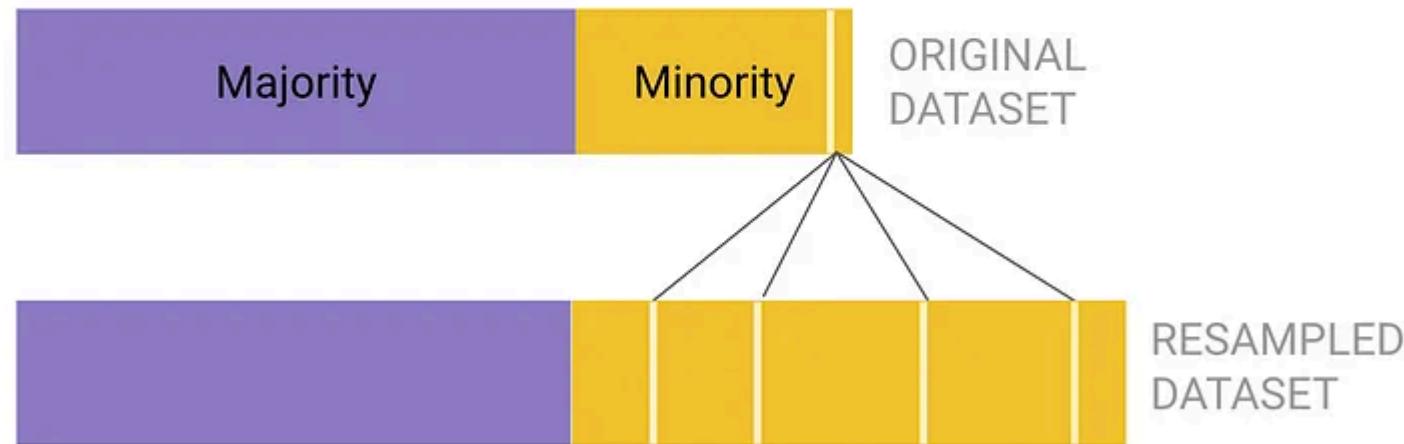


- Advantage: it does not introduce duplicates and/or artificial instances.
- Disadvantages:
 - ➡ Not all data points are used.
 - ➡ Potentially removing useful information.
- Better choice for data with very high class imbalance.

Up-sampling

Oversample minority class

TRAVIS TANG



Travis Tang from Medium

- Random up-sampling
- Disadvantages:
 - ➡ creates duplicated and/or artificial instances
 - ➡ can introduce bias and/or noise to the original data

Create synthetic samples of the minority class

- Synthetic Minority Over-sampling Technique (SMOTE) is a popular algorithm
- It creates synthetic samples from the minority class by:
 - ➡ Finding the k -nearest-neighbours for minority class observations
 - ➡ Randomly choosing one of the k -nearest-neighbours, then using it to create a similar but random new observation
- Be careful that you split your data into training/test sets before doing any oversample/SMOTE. Otherwise, you will leak information from training to test data set.

Handelling Class Imbalance via R

- Balanced Resampling with `ROSE:::ovun.sample()` can perform **under/over-sampling** or **hybrid** combines over- and under-sampling for balanced output.
- `themis:::step_smote` helps balance binary or multiclass classification problems by generating synthetic examples (via ***knn***) of the minority class.

