

STAT5003

Week 9 : Tree and Ensemble methods

Jaslene Lin

The University of Sydney



THE UNIVERSITY OF
SYDNEY

Readings and R functions covered


⚠ Important

- **Introduction to Statistical Learning**
 - ➡ Tree methods covered in Chapter 8
- **R** functions
 - ➡ `tree::tree` (Create Regression and Decision trees)
 - ➡ `base::sample` (Bagging)
 - ➡ `ranger::ranger` (construct random forests)
 - ➡ `gbm::gbm` (Gradient boosting machines)

This presentation is based on the [SOLES reveal.js Quarto template](#) and is licensed under a [Creative Commons Attribution 4.0 International License](#).

Decision Tree for Regression



- Baseball player salary data: how would you stratify it?
➡ **Salary** is colour-coded (NA values coded grey).



Decision Tree for Regression

- Overall, the tree segments the players into three regions of predictor space:


- ⇒ $R_1 = \{X \mid \text{Years} < 4.5\}$
- ⇒ $R_2 = \{X \mid \text{Years} \geq 4.5, \text{Hits} < 117.5\}$
- ⇒ $R_3 = \{X \mid \text{Years} \geq 4.5, \text{Hits} \geq 117.5\}$



Terminology for trees

- In keeping with the tree analogy:
 - ➡ The regions R_1 , R_2 , and R_3 are known as terminal nodes.
 - ➡ Decision trees are typically drawn upside down, in the sense that the leaves are at the bottom of the tree.
 - ➡ The points along the tree where the predictor space is split are referred to as internal nodes.
- In the Baseball player salary tree, the two internal nodes are indicated by **Years < 4.5** and **Hits < 117.5**.

Interpretation of Results



- **Years** is the most important factor in determining **Salary**:
 - ➡ Players with less experience earn lower salaries than more experienced players.
- Given that a player is less experienced:
 - ➡ The number of **Hits** that he made in the previous year seems to play little role in his **Salary**.
- Among players who have been in the major leagues for five or more years:
 - ➡ The number of Hits made in the previous year does affect Salary.
 - ➡ Players who made more Hits last year tend to have higher salaries.
- Obviously an over-simplification:
 - ➡ Compared to some other classification models (such as a regression model), it is easy to display, interpret, and explain.
 - ➡ Even easier than linear regression; you can interpret the decision tree to your grandparents

Details on tree building process

- In theory, the decision regions could have any shape.
- However, we choose to divide the predictor space into **high-dimensional rectangles, or boxes**, for simplicity and for ease of interpretation of the resulting predictive model
- The goal is to find boxes R_1, \dots, R_J that minimizes the residual sum of squares (RSS), given by:

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

➡ \hat{y}_{R_j} is the mean response for the training observations within the j^{th} box

Tree Building via Recursive Binary Splitting

- We want to partition the feature space into J distinct regions.
- **But** trying every possible partition is computationally **infeasible**.
- Instead, we use a **top-down greedy approach (recursive binary splitting)**:
 - ➡ At each step, choose the best split.
 - ➡ Recursively repeat the process on resulting subspaces.

Recursive Binary Splitting: Key Ideas

- **Top-down**: Start at the root node (entire dataset) and split down.
- **Greedy**: At each step, choose the split that most reduces RSS or impurity *at that step*.
- **Recursive**: Once a split is made, repeat the process on each child node.


Note

This method does **not** look ahead – it optimises **locally**, not globally.

Recursive Splitting Visual Example (One-Dimensional Simulated Data)


Recursive Binary Splitting (Toy Example)

Dashed red lines show sequential splits on predictor x



Overfitting with Recursive Binary Splitting

- Trees are grown via **recursive binary splitting** can lead to **deep, complex trees** that fit training data perfectly.
- Such trees may **overfit** as it captures **noise** instead of signal
- This leads to **high variance** and **poor generalisation**



```
1 # Step 1: Predict on the training data
2 predicted <- predict(tree.full, newdata = test_data)
3
4 # Step 2: Extract the actual values
5 actual <- test_data$Salary
6
7 # Step 3: Compute the MSE
8 mse <- mean((actual - predicted)^2)
9 mse
```

[1] 162111.2

Tree Pruning

Note

Tree pruning is like regularisation – we sacrifice training accuracy to gain prediction stability (lowering the variance).

Cost-Complexity Pruning (the weakest link pruning)

$$C_\alpha(T) = \sum_{m=1}^{|T|} \sum_{i:x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha|T|$$


T : subtree ; R_m : region (leaf) m

\hat{y}_{R_m} : mean response in region m


α : tuning parameter controls a trade-off between the subtree's complexity and its fit to the training data.

Cost-Complexity Pruning (the weakest link pruning)

► Code



► Code




► Code

```
[1] 167258.6
```

Using Decision Trees for prediction

- We divide the predictor space:
 - ➡ That is, the set of possible values for X_1, X_2, \dots, X_p , into J distinct and non-overlapping regions, R_1, R_2, \dots, R_J .
- For every observation that falls into the region R_j , we make the same prediction, which is simply the mean of the response values for the training observations in R_j .



Decision Trees for Classification

- Very similar to a regression tree
 - ➡ Exception being the prediction of a qualitative response rather than a quantitative one.
- For a classification tree
 - ➡ Inspect the region that the observation belongs and predict the most commonly occurring class in that region.

Gini index

- Just as in the regression setting, we use recursive binary splitting to grow a classification tree.
- In the classification setting, RSS cannot be used as a criterion for making the binary splits.
- Alternative measure such as Gini index is used instead.
- The Gini index is defined by

$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk}).$$

➡ \hat{p}_{mk} represents the proportion of training observations in the m^{th} region that are from the k^{th} class.

Gini index


- The Gini index is a measure of total variance across the K classes.
 - ➡ It takes on a small value if all of the \hat{p}_{mk} values are close to zero or one.
 - ➡ This occurs when there is a clear majority class!
- For this reason the Gini index is referred to as a measure of node **purity**.
 - ➡ A small value indicates that a node contains predominantly observations from a single class.

Entropy





- An alternative to the Gini index is entropy, given by

$$D = - \sum_{j=1}^J \sum_{k=1}^K \hat{p}_{jk} \log \hat{p}_{jk}$$

- It turns out that the Gini index and the Entropy are very similar numerically



Linear Model versus Tree



- If the relationship between the features and the response is well approximated by a linear model as shown in the top panel, then fitting a linear model (Graph 1) outperforms the tree-based model (Graph 2).
- If the relationship is complex and non-linear as shown on the bottom panel, then the tree-based model (Graph 4) is more appropriate.

Advantages and disadvantages of trees




Advantages:

- Trees are very easy to explain to people
- Some people believe that decision trees closely relate to human decision-making
- Trees can be displayed graphically
- Can handle different data types and doesn't require scaling

Disadvantages:

- Trees do not have the same level of predictive accuracy as some of the other regression and classification approaches we have discussed

A single decision tree is prone to over-fitting



Ensemble of trees

- Build **multiple** decision trees
- Combine their results to form a more **stable and accurate** predictor

Key Ideas

- Take **samples** of trees (e.g. using bootstrapping)
- Use **random subsets of features** at each split
- **Aggregate** predictions:
 - ➡ **Majority vote** (classification)
 - ➡ **Average** (regression)

Popular Methods

- **Bagging**: Build trees in parallel and average results
- **Random Forest**: Bagging + random feature selection
- **Boosting**: Build trees sequentially to improve errors

The strength of ensemble methods lies in **reducing variance** and **improving robustness** through aggregation.

"Many weak learners can form a strong learner."

Bagging (Bootstrap Aggregating)


- Bootstrap aggregating, or bagging, is a general purpose procedure for reducing the variance of a statistical learning method.
 - ➡ It is particularly useful and frequently used in the context of decision trees.
- Recall that given a set of n independent observations Z_1, \dots, Z_n each with a variance of σ^2 :
 - ➡ The variance of the mean \bar{Z} is given by σ^2/n .
- In other words, averaging a set of observations reduces variance.
- Since it is typically not possible to have access to multiple training sets:
 - ➡ We can use **bootstrapping** to create multiple training sets.

Bagging continued

- Use bootstrapping to take repeated samples from a (single) training data set.
- In this approach, we generate B different bootstrapped training data sets.
 - ➡ Train our method of the b^{th} bootstrapping set in order to obtain $\hat{f}_b^*(x)$, the prediction at a point x .
- Average all the observations to obtain:

$$\hat{f}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b^*(x).$$

- This is called bagging.



Out of bag error estimation


- It turns out that there is a very straightforward way to estimate the test error of a bagged model.
- Recall that the key to bagging is that trees are repeatedly fit to bootstrapped subsets of the observations.
 - ➡ One can show that on average, each bagged tree makes use of around $\mathbf{2/3}$ of the observations when the sample size n is large.
- The remaining $\mathbf{1/3}$ of the observations not used to fit a given bagged tree are referred to as the out-of-bag (OOB) observations.
- We can predict the response for the i^{th} observation using each of the trees in which that observation was OOB.
 - ➡ This will yield around $B/3$ predictions for the i^{th} observation, which we average.
 - ➡ This estimate is essentially the leave-one-out (LOO) cross-validation error for bagging, if B is large.

From bagging to Random Forest

- Random forests (sometimes) provide an improvement over bagged trees by way of a small tweak that decorrelates the trees. This reduces the variance when we average the trees.
- As in bagging, we build a number of decision trees on bootstrapped training samples.
- But when building these decision trees, each time a split in a tree is considered, a random selection of m predictors is chosen as split candidates from the full set of p predictors.
 - ➡ The split is allowed to use only one of those m predictors.
- A fresh selection of m predictors is taken at each split, and typically we choose $m \approx \sqrt{p}$.
 - ➡ That is, the number of predictors considered at each split is approximately equal to the square root of the total number of predictors.

Random Forest algorithm pseudocode

```
1 Set number of models to build, B
2 for i = 1 to B
3     Generate a bootstrap sample of the original data
4     Train a tree model on this sample where
5         for each split
6             Randomly select m (< p) of the original predictors
7             select the best predictor among the m predictors and partition the data
8     endfor
9 endfor
```



Ensemble methods via boosting

Boosting

- Like bagging, boosting is a general approach that can be applied to many statistical learning methods for regression or classification. We only discuss boosting for decision trees.
- Recall that bagging involves creating multiple copies of the original training data set using the bootstrap, fitting a separate decision tree to each copy, and then combining all of the trees in order to create a single predictive model.
- Notably, each tree is built on a bootstrap data set, **independent of other trees**.
- Boosting works in a similar way, except that the trees are grown *sequentially*.
 - ➡ Each tree is grown using information from previously grown trees.

Idea behind boosting

- Unlike fitting a single large decision tree to the data, which amounts to *fitting the data hard* and potentially overfitting, the boosting approach instead *learns slowly*.
- Each new model attempts to **correct the errors** made by the previous ones.
- The final model is a **weighted sum** of weak learners (usually decision trees).
- By fitting small trees to the residuals, we slowly improve $\hat{f}(x)$ in areas where it does not perform well.
- There is a shrinkage parameter λ that slows the process down even further, allowing more and different shaped trees to attack the residuals.

Boosting for regression trees

1. Set $\hat{f}(x) = \mathbf{0}$ and $r_i = y_i$ for all i in the training set
 - r_i denotes the i^{th} residual, and y_i is the outcome
2. For $b = 1, 2, \dots, B$:
 - Fit a tree \hat{f}_b with d splits ($d + 1$ terminal nodes) to the new training data (X, r)
 - ⇒ That is, r is the new response value
 - Update \hat{f} by adding in a shrunken version of the new tree
 - ⇒ $\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}_b(x)$
 - Update the residuals
 - ⇒ $r_i \leftarrow r_i - \lambda \hat{f}_b(x)$
3. Compute the final model

$$\hat{f}(x) = \sum_{i=1}^B \lambda \hat{f}_b(x)$$

Parameters to tune in boosting

- Number of trees B
- The shrinkage parameter λ : a small positive number
 - ➡ Typical values are between 0.01 and 0.001
- Number of split d in each tree
 - ➡ Controls the complexity of the boosted ensemble
 - ➡ If $d = 1$, then the tree is just a stump. This actually usually works quite well

Gradient Boosting in Action (R Demo)


► Code

```
CV: 1  
CV: 2  
CV: 3  
CV: 4  
CV: 5
```

► Code

► Code

```
[1] 223
```




Gradient Boosting in Action (R Demo)

► Code

```
CV: 1  
CV: 2  
CV: 3  
CV: 4  
CV: 5
```

► Code

Boosted Model Fit




Other boosting algorithms

- AdaBoost
- Stochastic gradient boosting
- XGBoost


AdaBoost

- Short for Adaptive Boosting, one of the first boosting algorithms for classification
- Convert a set of weak classifiers (*decision stumps*) into a strong one
- Basic idea
 - ➡ At each iteration, reweight the data to place more weight on data points that the classifier got wrong
- At the end, combine all the weak classifiers by taking a weighted combination
 - ➡ Put more weight on the weak classifiers with the higher accuracies


AdaBoost plot



Iteration 1




Iteration 2




Iteration 3

Final Model


$$f(x) = 0.46$$



$$+ 0.29$$



$$+ 0.46$$




Stochastic gradient boosting

- Use the idea behind bagging
- In each iteration of stochastic boosting, a sample of the training set instead of the full training set is used
- Instead of a bootstrap sample (with replacement), the algorithms samples a fraction of the training set
- This introduced randomness can improve performance
- Paper by [Stochastic gradient boosting](#) covers this technique

Stochastic gradient boosting plot

$N = 500$



- Taken from Stochastic gradient boosting

XGBoost

- XGBoost is short for eXtreme Gradient Boosting
- It is a library for high performance gradient boosting models written in C++ with implementations in Python, R, and Julia
- Designed for speed: up to 10 times faster than the `gbm` package
- Has been very popular in recent years and has won a number of machine learning competitions (especially on tabular data)
- Supports regularization
- Can handle missing data

Bagging vs Boosting

- Both ensemble methods get N learners from 1 learner
 - ➡ Built independently for bagging
 - ➡ Built sequentially for boosting
- Trees built in boosting are weak learners (sometimes just a stump) while trees in Random Forest have higher complexity
- Fewer parameters to tune in Random Forest, many more in Boosting (depending on which variation you are using)
- Both combine outputs from N trees

Summary

- Decision trees are simple and interpretable models for regression and classification.
- However, they are often not competitive with other methods in terms of prediction accuracy.
- Bagging, random forests, and boosting are good methods for improving the prediction accuracy of trees.
 - ➡ They work by growing many trees on the training data and then combining the predictions of the resulting ensemble of trees.
- The latter two methods, random forests and boosting, are among the state-of-the-art methods for supervised learning.
 - ➡ Their results can be difficult to interpret.
 - ➡ But we can still get **relative importance** of each predictor.
 - ➡ For classification, the relative importance is computed using the mean decrease in Gini index and expressed relative to the maximum.