

STAT5003

Week 10 : Monte Carlo

Jaslene Lin

The University of Sydney



THE UNIVERSITY OF
SYDNEY

Readings

! Important

- Easily accessible text in [Explorations in Monte Carlo Methods](#)

This presentation is based on the [SOLES reveal.js Quarto template](#) and is licensed under a [Creative Commons Attribution 4.0 International License](#).

Monte Carlo Methods

What are Monte Carlo methods?

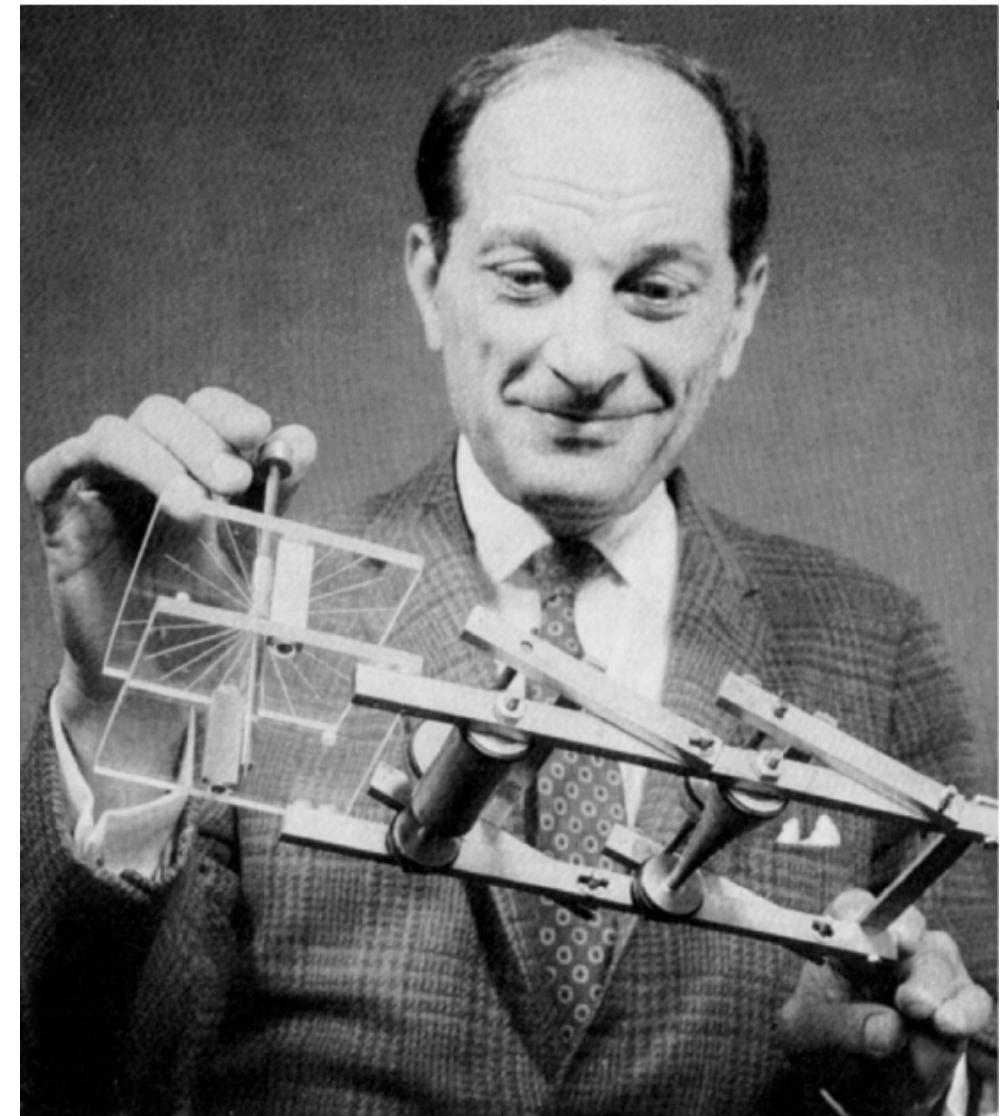
- Monte Carlo are a class of computational methods that can be applied to a wide range of problems
- Key aspect of Monte Carlo methods are that they rely on **random sampling**
- Generally provide approximate solutions
 - ➡ Not exact solutions
- Used in cases where analytic solutions don't exist or are too difficult to implement
- Monte Carlo methods are sometimes referred to as stochastic simulation

History of the Monte Carlo method

- Invented by a Polish mathematician called Stanislaw Ulam in the late 1940s
- Inspiration during recovery from illness with a thought experiment:

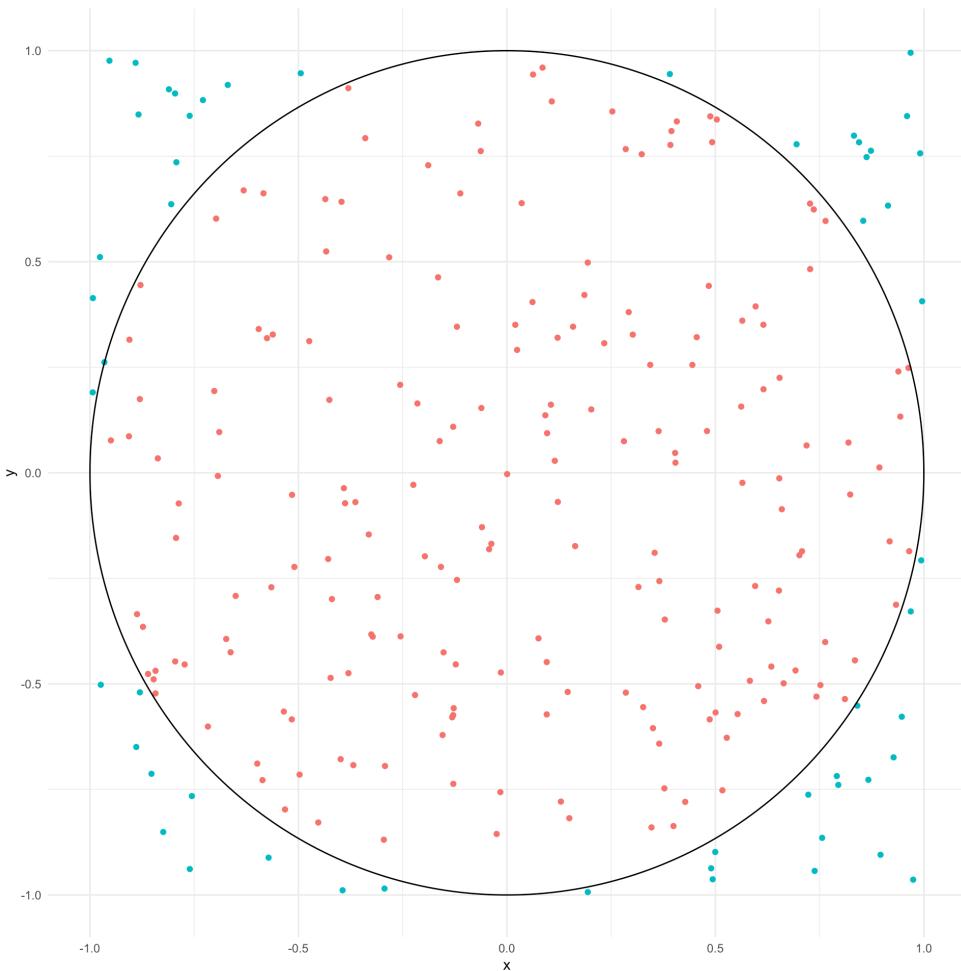
The question was what are the chances that a Canfield Solitaire laid out with 52 cards will come out successfully? After spending a lot of time trying to estimate them by pure combinatorial calculations, I wondered whether a more practical method than "abstract thinking" might not be to lay it out say one hundred times and simply observe and count the number of successful plays.

- Monte Carlo was the code name given for the project
 - ➡ Inspired by the Monte Carlo casino in Monaco where Ulam's uncle would borrow family money to gamble



By Originally uploaded by Deer*lake (Transferred by Deer*lake) - Originally uploaded on en.wikipedia, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=18403886>

Monte Carlo estimation of π



- Area of circle: πr^2 ; area of square: $4r^2$
⇒ r is the radius
- A Monte Carlo method to estimate π
 1. Randomly draw N points within unit square at random
 2. Count the R points which are inside the unit circle
 3. Compute the ratio R/N and estimate π as $4R/N$

Using Monte Carlo to estimate π

In this example, the mathematical statistics of the procedure is:

- Write the parameter we want to estimate, i.e., π , as an expectation
- Represent the parameter as a sample approximation
- Let $f(\cdot)$ be the probability density function (pdf) of X and A be the support of X

$$\mathbb{E}[X] = \int_A t f(t) dt \approx \frac{1}{N} \sum_{i=1}^N X_i$$

- The law of large numbers guarantees that as the number of samples we draw increase, the estimator converges to the true parameter value almost surely

Expectation of a random variable

From last slide, we can write the expected value of a random variable \mathbf{X} as a sum:

$$\mathbb{E}[\mathbf{X}] = \int_A t \cdot f(t) dt \approx \frac{1}{N} \sum_{i=1}^N X_i$$

We can replace \mathbf{X} with a function $\mathbf{g}(\mathbf{X})$. Then the previous equation becomes:

$$\mathbb{E}[\mathbf{g}(\mathbf{X})] = \int_A g(t) \cdot f(t) dt \approx \frac{1}{N} \sum_{i=1}^N g(X_i)$$

Here, we assume to draw $X_i \sim f$

Monte Carlo Integration example

Goal is to evaluate the following definite integral

$$\int_0^1 e^{-x^2/2} dx \approx \frac{1}{N} \sum_{i=1}^N g(X_i).$$

Find a random variable such that the above holds.

Need to be careful such that:

- Domain of the random variable and g agree with the above definition

- One example is:

$$\Rightarrow g(x) = e^{-x^2/2}$$

$\Rightarrow X_i$ is a uniform random variable on $(0, 1)$

$$\Rightarrow \text{Then } \mathbb{E}[g(X)] = \int_0^1 e^{-x^2/2} dx \approx \frac{1}{N} \sum_{i=1}^n g(X_i)$$

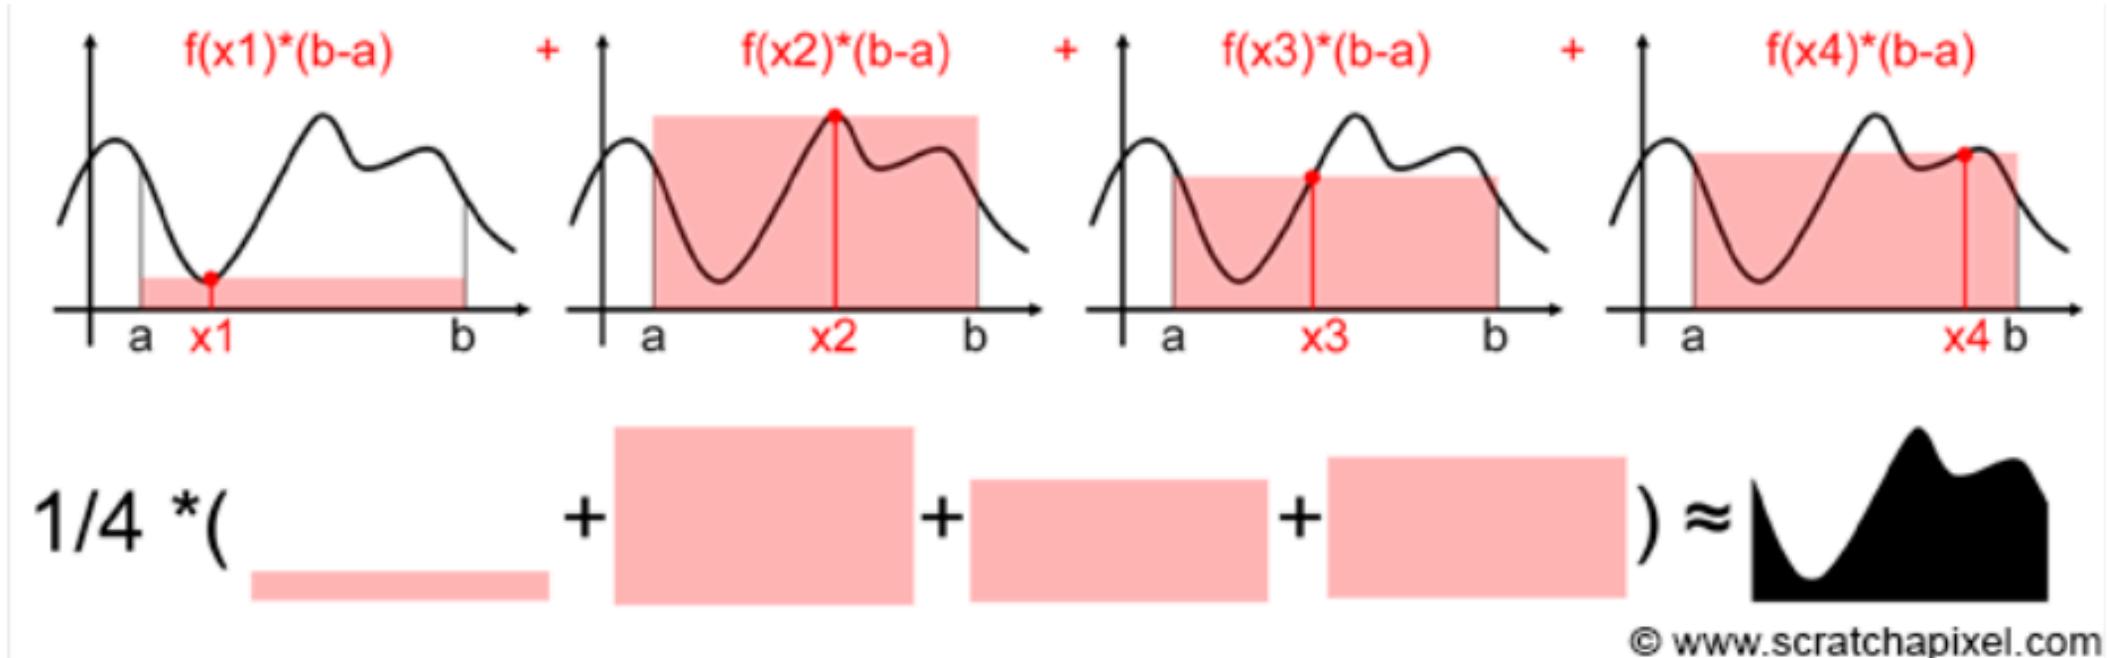
```
1 ## define the function
2 f <- function(x){
3   exp(-x^2/2)
4 }
5
6 ## exact solution
7 integrate(f, lower = 0, upper = 1)
```

0.8556244 with absolute error < 9.5e-15

```
1 ## Monte Carlo estimation
2 set.seed(5003)
3 x <- runif(1000, 0, 1)
4 mean(f(x))
```

[1] 0.8545841

Monte Carlo visually



Simulating random variable

- Let's say we know how to simulate random variables that has a **uniform** distribution on the interval **[0, 1]**
 - Call this uniformly distributed random variable **U**
 - Can do that in R with **runif**

How do we simulate random variables that can have any probability distribution?

In R, we can draw random Gaussian variables using the function **rnorm** but how do these functions really work?

Two methods (others exist):

1. Inverse-transform method
2. Acceptance-rejection method

Inverse transform method

- Let's say X is a random variable that has a cumulative distribution function (cdf) F
- Remember that cdf is the integral of the pdf

$$F(x) = \int_{-\infty}^x f(t) dt$$

- Recall U is a uniformly distributed random variable over $[0, 1]$
- If the inverse of $F(x)$ exists, then we can generate X as:

$$X = F^{-1}(U)$$

Example: Exponential distribution

The exponential distribution has a pdf of the form:

$$f(x) = \begin{cases} \lambda e^{-\lambda x} & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

and a cdf of the form:

$$F(x) = \begin{cases} 1 - e^{-\lambda x} & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

The inverse cdf is:

$$F^{-1}(p) = -\frac{\log(1-p)}{\lambda}, \quad 0 < p < 1$$

To sample from the exponential distribution, we can do the following:

1. Sample $U \sim \text{Uniform}(0, 1)$

2. Set $X = -\frac{\log(1-U)}{\lambda}$

Acceptance-Rejection method

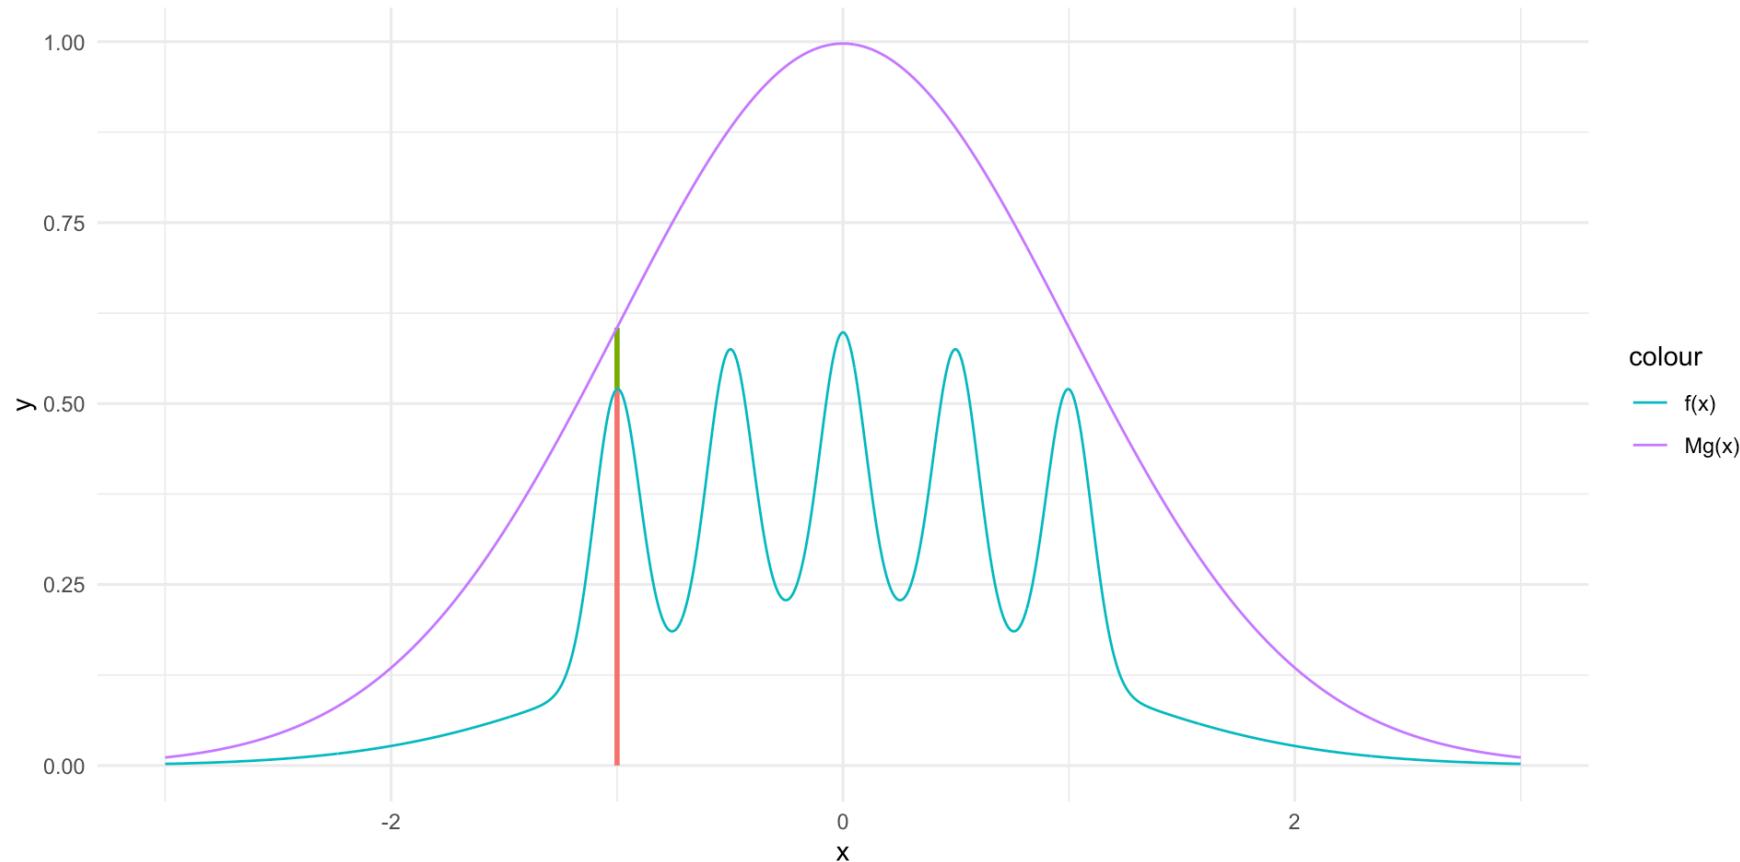
- The problem with the inverse transform method is that you need to know the functional form of the inverse of the cdf
- Acceptance-rejection method is a more general method can simulate “difficult” distributions

Acceptance-Rejection method

- Given two probability densities $f(x)$ and $g(x)$, with $f(x) < Mg(x)$ for all x
- If $g(x)$ is a distribution that can be easily sampled, then we can sample $f(x)$ using the following procedure:
 1. Draw a random variable $X \sim g(x)$
 2. Accept X with probability $\frac{f(x)}{Mg(x)}$
 3. Repeat steps 1 and 2 until you have the desired number of random samples

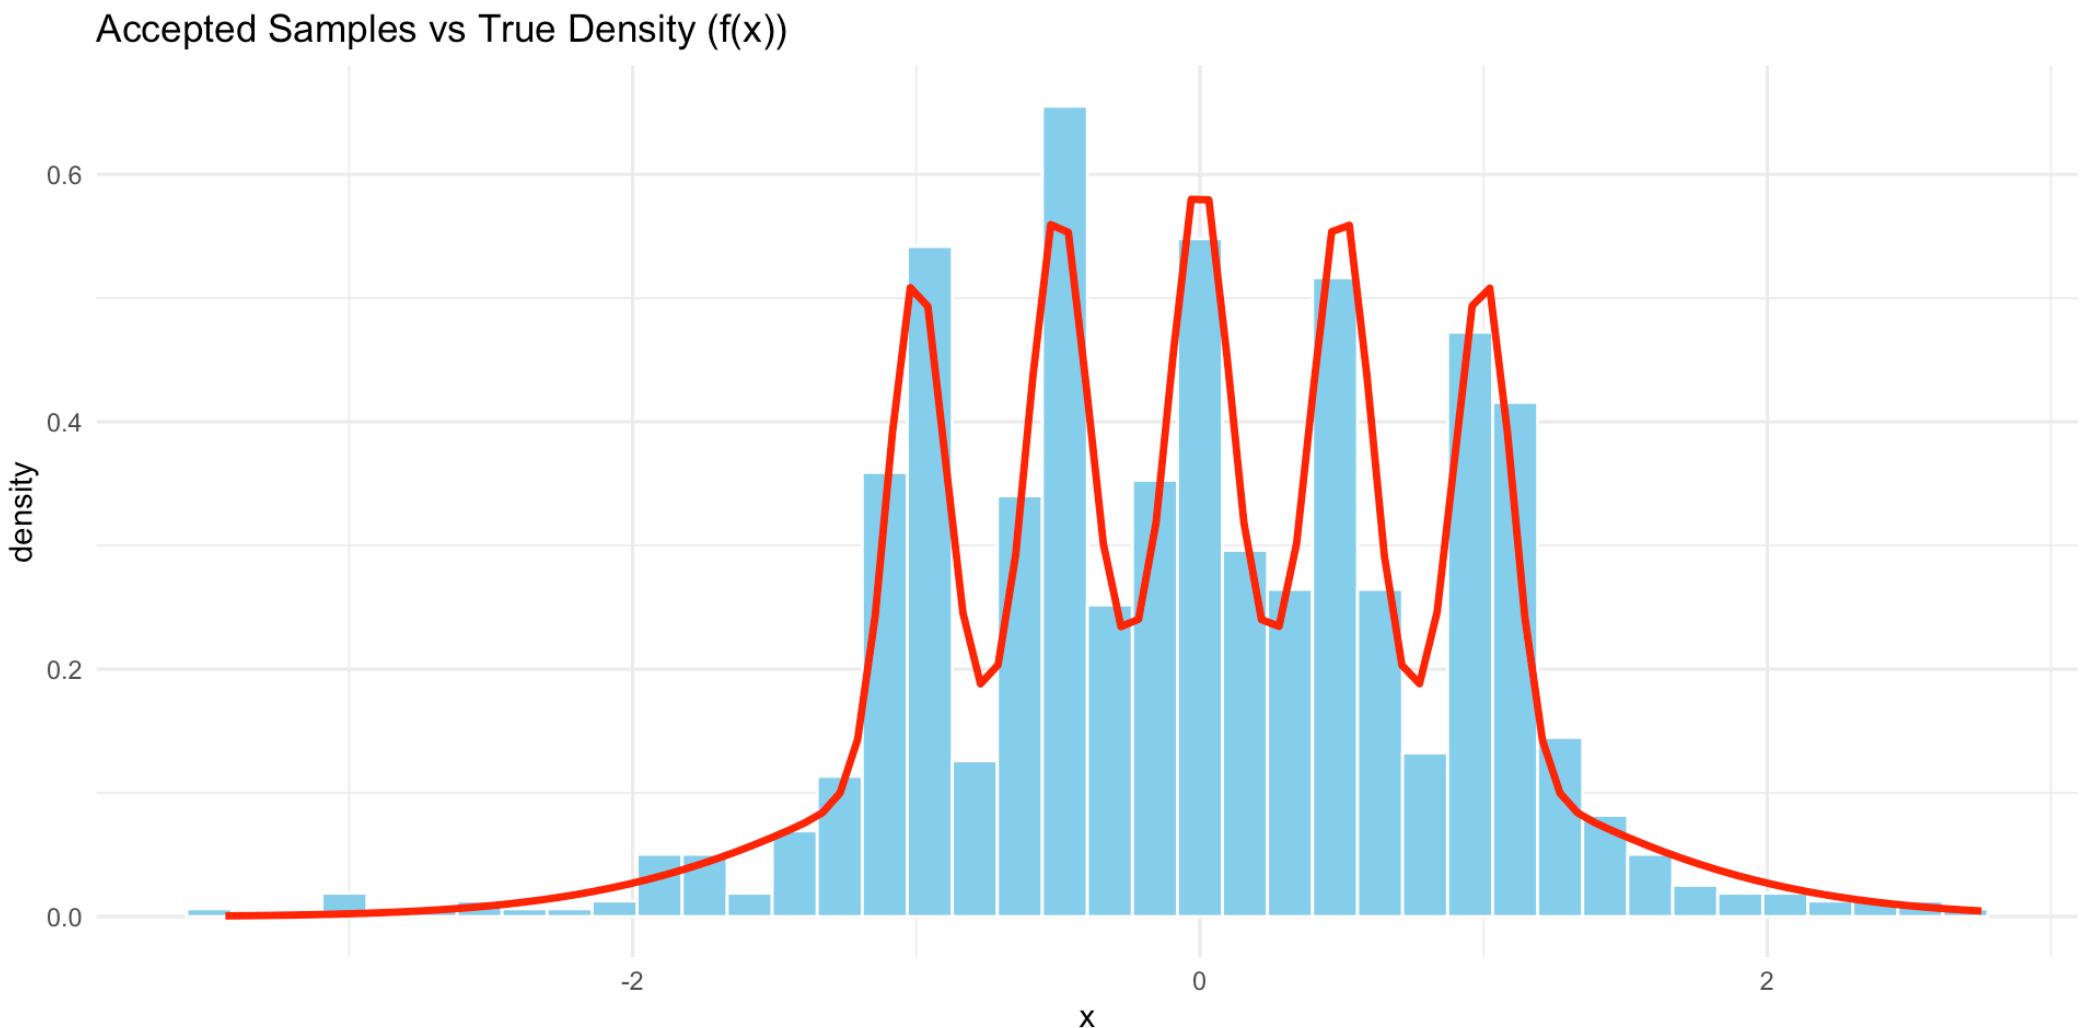
Acceptance-Rejection plot

- $f(x)$ is a normal mixture distribution, difficult to sample directly
- $M = 2.5$
- $g(x) \sim N(0, 1)$



Acceptance-Rejection R Demo

► Code



Problems with the Acceptance-rejection method

- The envelop distribution $g(x)$ needs to closely resemble $f(x)$ for this method to work well
- If $g(x)$ does not match $f(x)$ well:
 - ➡ The method will draw a lot of unwanted samples hence is not very efficient
 - ➡ Become very problematic when the data are high-dimensional
 - ➡ Even at dimensions of ~ 10 , this method becomes very inefficient, i.e., you need to draw lots of samples before one sample is accepted

Monte Carlo simulation examples

Pop Mart Blind Box

ALL MEMBERS



- Each Pop Mart Monster blind box costs **\$22**.
- There are **10 unique figures** in the full set.
- Each box contains **one random figure**.

? How many boxes would you need to purchase – and how much money would you expect to spend – to collect **all 10 figures**?

Analytical Solution (before we see the Monte Carlo)

Expected number of shops you need to do:

$$\mathbb{E}[T] = n \left(\frac{1}{1} + \frac{1}{2} + \cdots + \frac{1}{n} \right) \approx n \log n + 0.5772n + 0.5$$

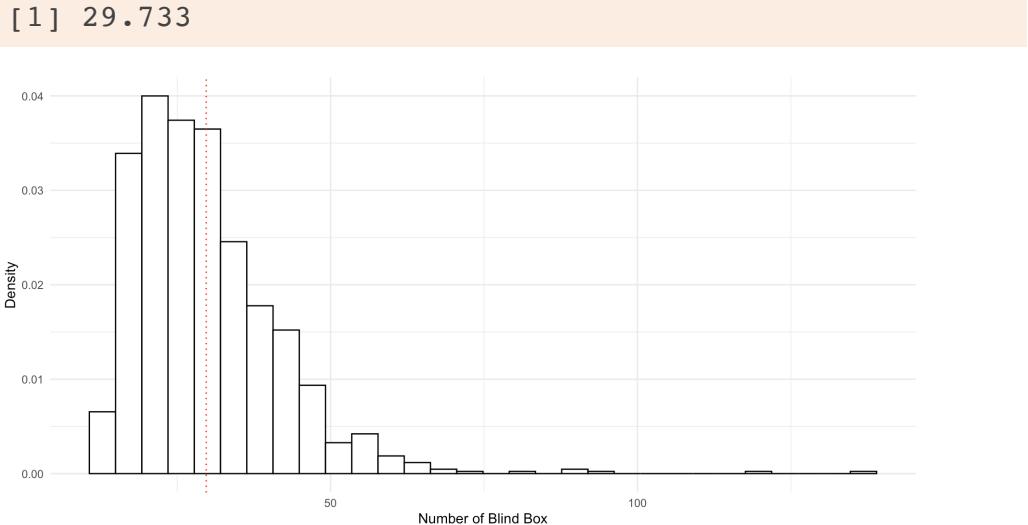
- T is the number of shops you need to do, and n is the number of items to collect

For the example, we have $n = 10$ (the number of unique figures) so $\mathbb{E}[T] \approx 30$

- On average, you need to buy about 30 boxes.
- The expected total money spent is $30 \times 22 = 660$ dollars

Monte Carlo Simulation

```
1 n.shops.sim <- c()
2 for(i in 1:1000) {
3   collected <- c(); n.shops <- 0
4   while(length(collected) < 10) {
5     newitem <- sample(10, 1)
6     collected <- union(collected, newitem)
7     n.shops <- n.shops + 1
8   }
9   n.shops.sim <- c(n.shops.sim, n.shops)
10 }
11 mean.sim <- mean(n.shops.sim)
12 mean.sim ## the average number of eligible shops.
13 ggplot(data.frame(x = n.shops.sim)) +
14   theme_minimal() +
15   geom_histogram(
16     aes(x = x, y = after_stat(density)),
17     colour = "black", fill = "white") +
18   labs(x = "Number of Blind Box", y = "Density") +
19   geom_vline(xintercept = mean.sim,
20             linetype = "dotted",
21             colour = "red")
```



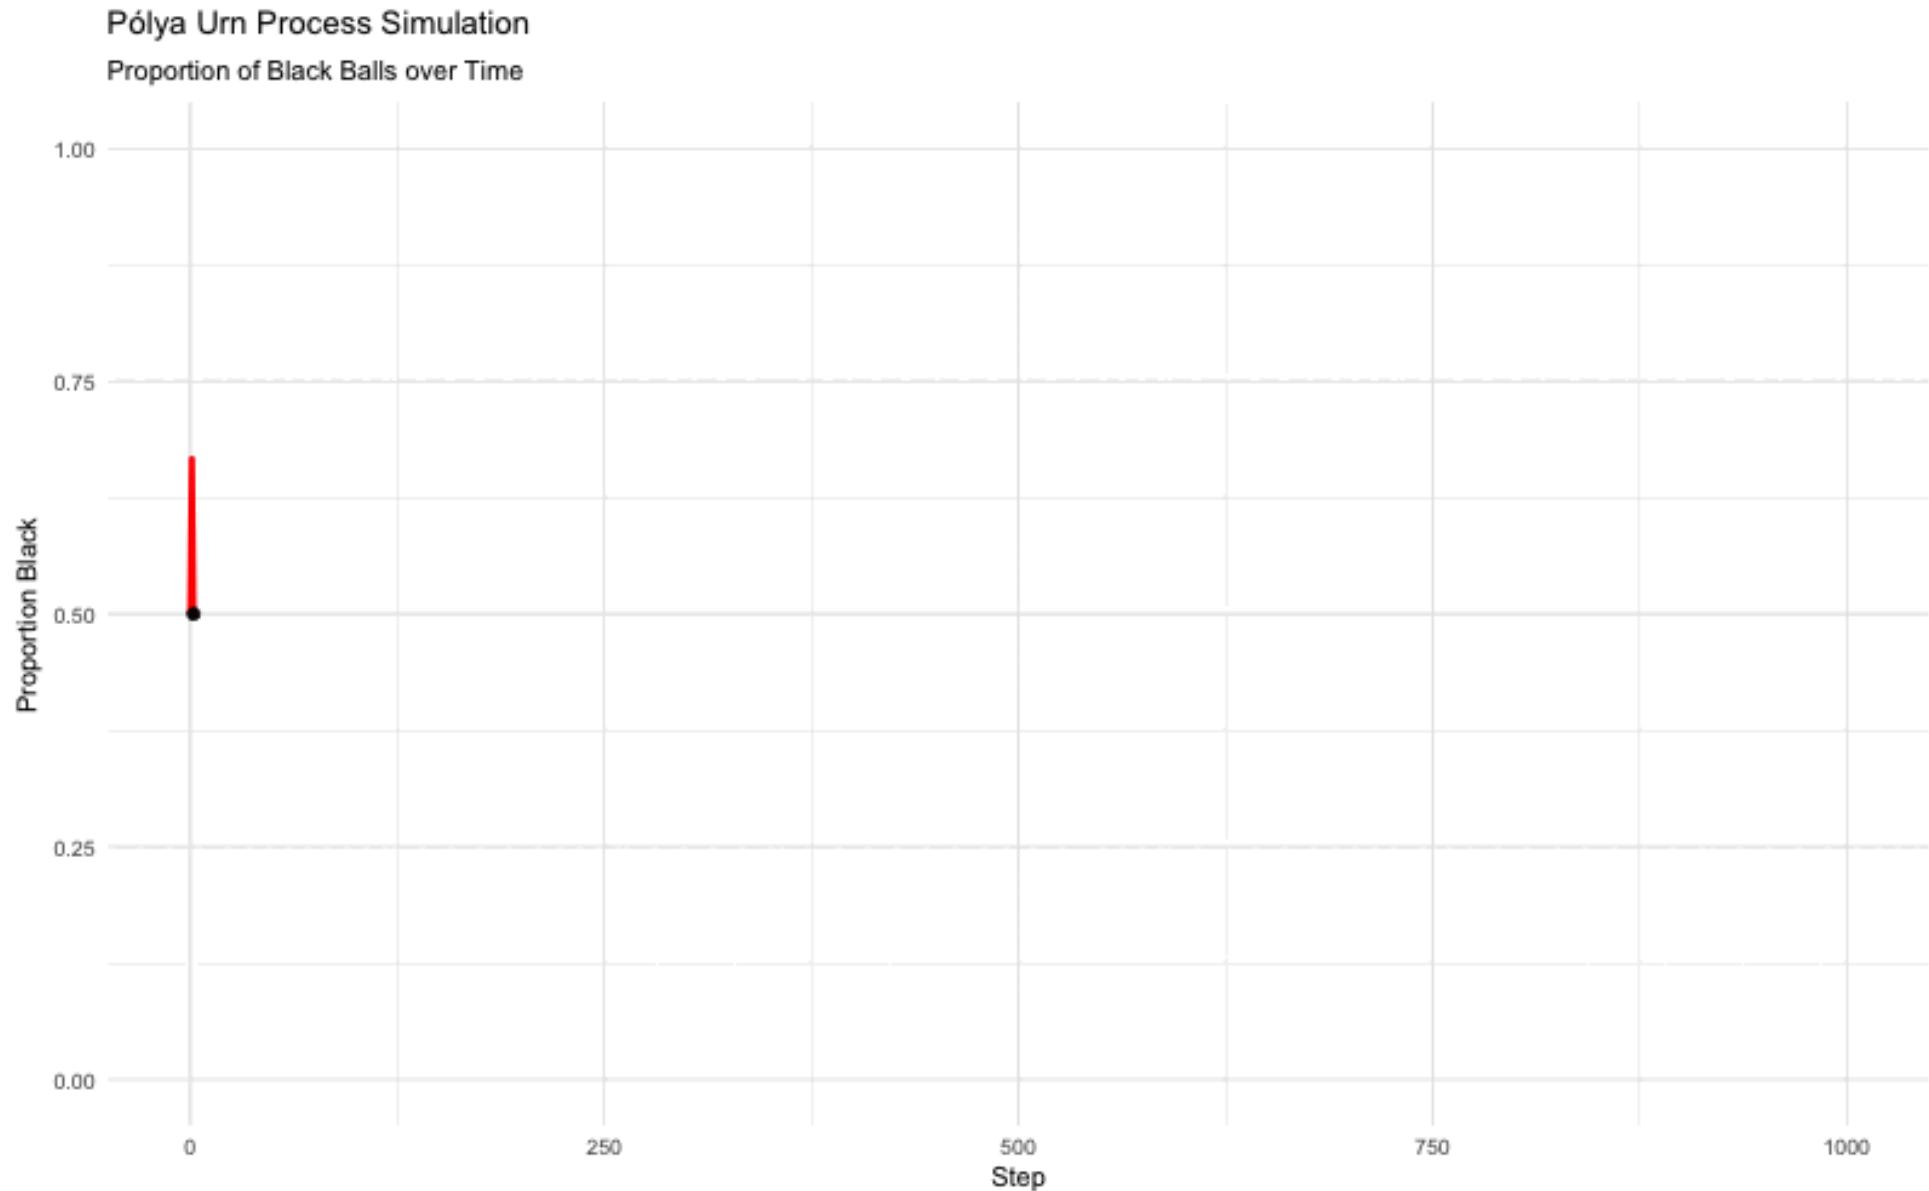
Polya urn problem

- Let an urn start with 1 black ball and 1 white ball
- Repeatedly draw a single ball from the urn
- Each time a ball is drawn:
 - ➡ You put back that ball plus one extra ball of the same colour back into the urn

After 1000 draws, how many black balls do you expect to be in the urn?

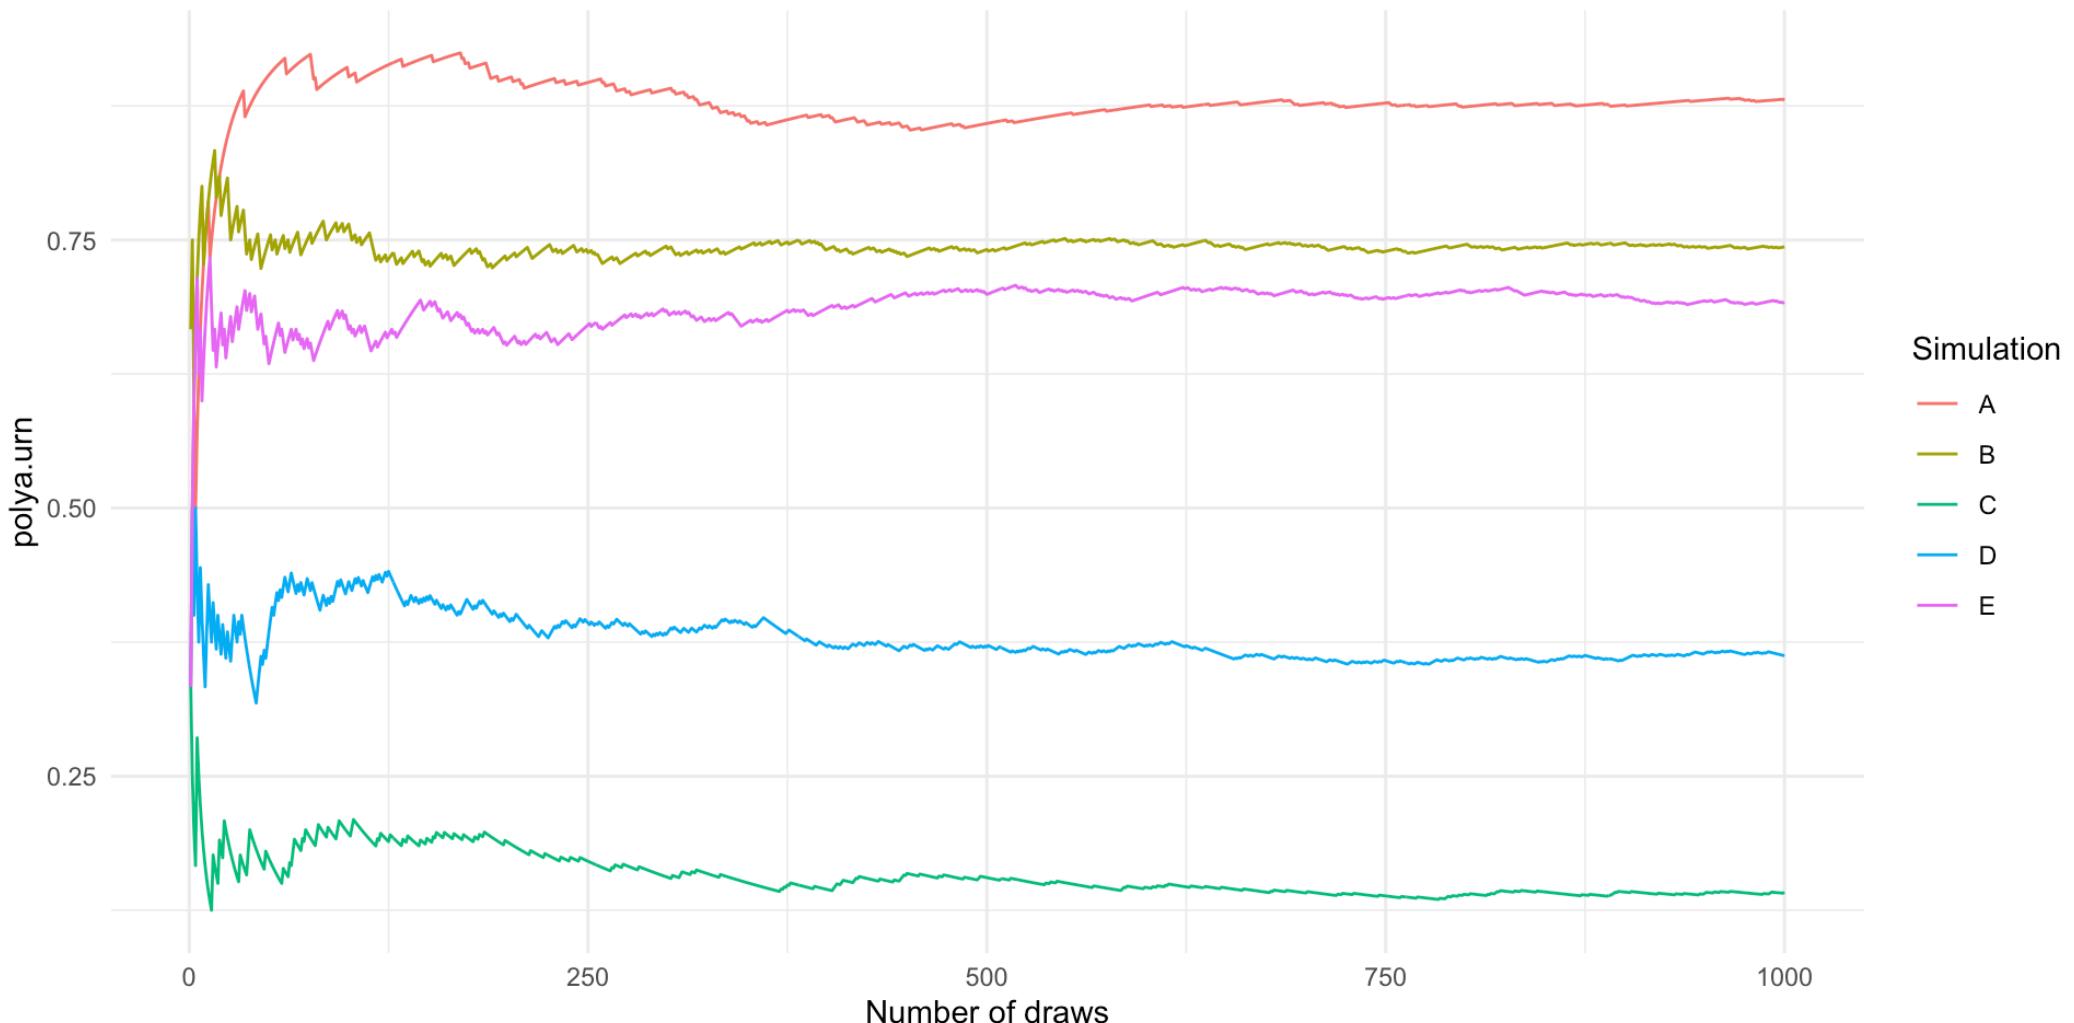
- This is an example of using Monte Carlo to simulate a process

Monte Carlo simulation of the Polya urn problem

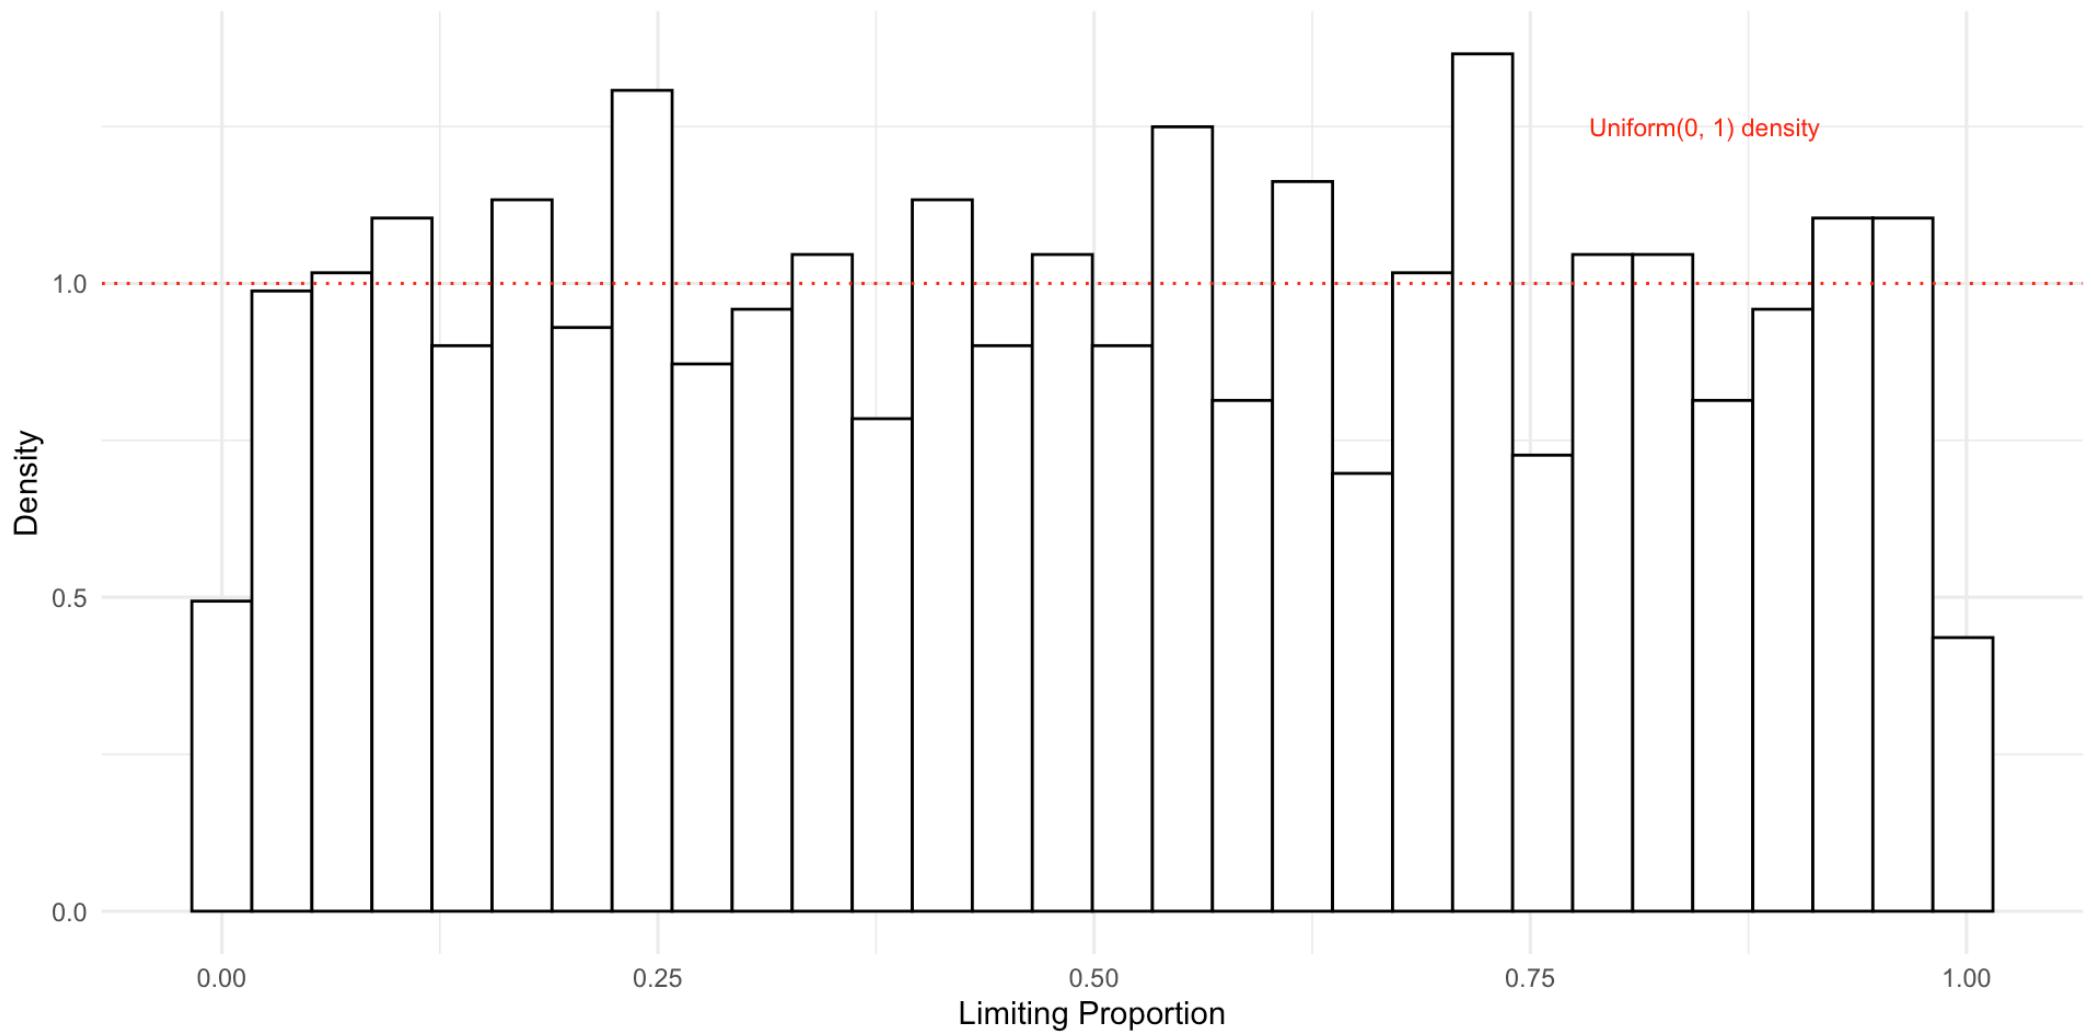


Monte Carlo simulation of the Polya urn problem

► Code



Asymptotic distribution of the proportion of black balls



- In the limit, the distribution is $U(0, 1)$

Call Option Pricing

An European call option gives you the right to buy a stock at a particular price at expiry

Example: 90 day call option of a stock with \$105 strike. Let's say a stock is currently priced at \$100. If it hits \$108 at day 90, then you can exercise the option and make a \$3 profit from this stock. If its price is below \$105 at day 90, then you make nothing.

- Question: **How much should you pay for this option today?**

Think of this as the **insurance premium** you'd pay to guarantee the right to buy later at a known price.

A put option gives you the right to sell a stock at a particular price at expiry

Option Payoff at Expiry

- At expiration, the **payoff** of the call option is:

$$\text{Payoff} = \max(S_T - K, 0)$$

where S_T is the stock price at expiry and K is the strike price.

- There are three possible scenarios:
 - ➡ $S_T < K$: Payoff = 0 (option is worthless)
 - ➡ $S_T = K$: Payoff = 0
 - ➡ $S_T > K$: Payoff = $S_T - K$

Simulate stock price as a random (stochastic) model

Equation for the future price S of a stock:

$$dS = S(\mu dT + \sigma \sqrt{dT} \mathcal{N})$$

- μ is the growth rate of the stock price
- σ is known as the volatility
 - ➡ Think of this as the variance of the stock price
- \mathcal{N} is a standard normal random variable
- dT is one unit of time

Monte Carlo Simulation

1. **Simulate** many possible future stock prices S_T based on the above equation.
2. For each simulation, **calculate payoff**.
3. The price you are willing to pay is the expected value of the payoff $E(\text{Payoff})$.

Examples of a stock price

Example: 90 day call option of a stock with \$105 strike. currently priced at \$100.

► Code

```
[1] 8.723601
```

► Code

